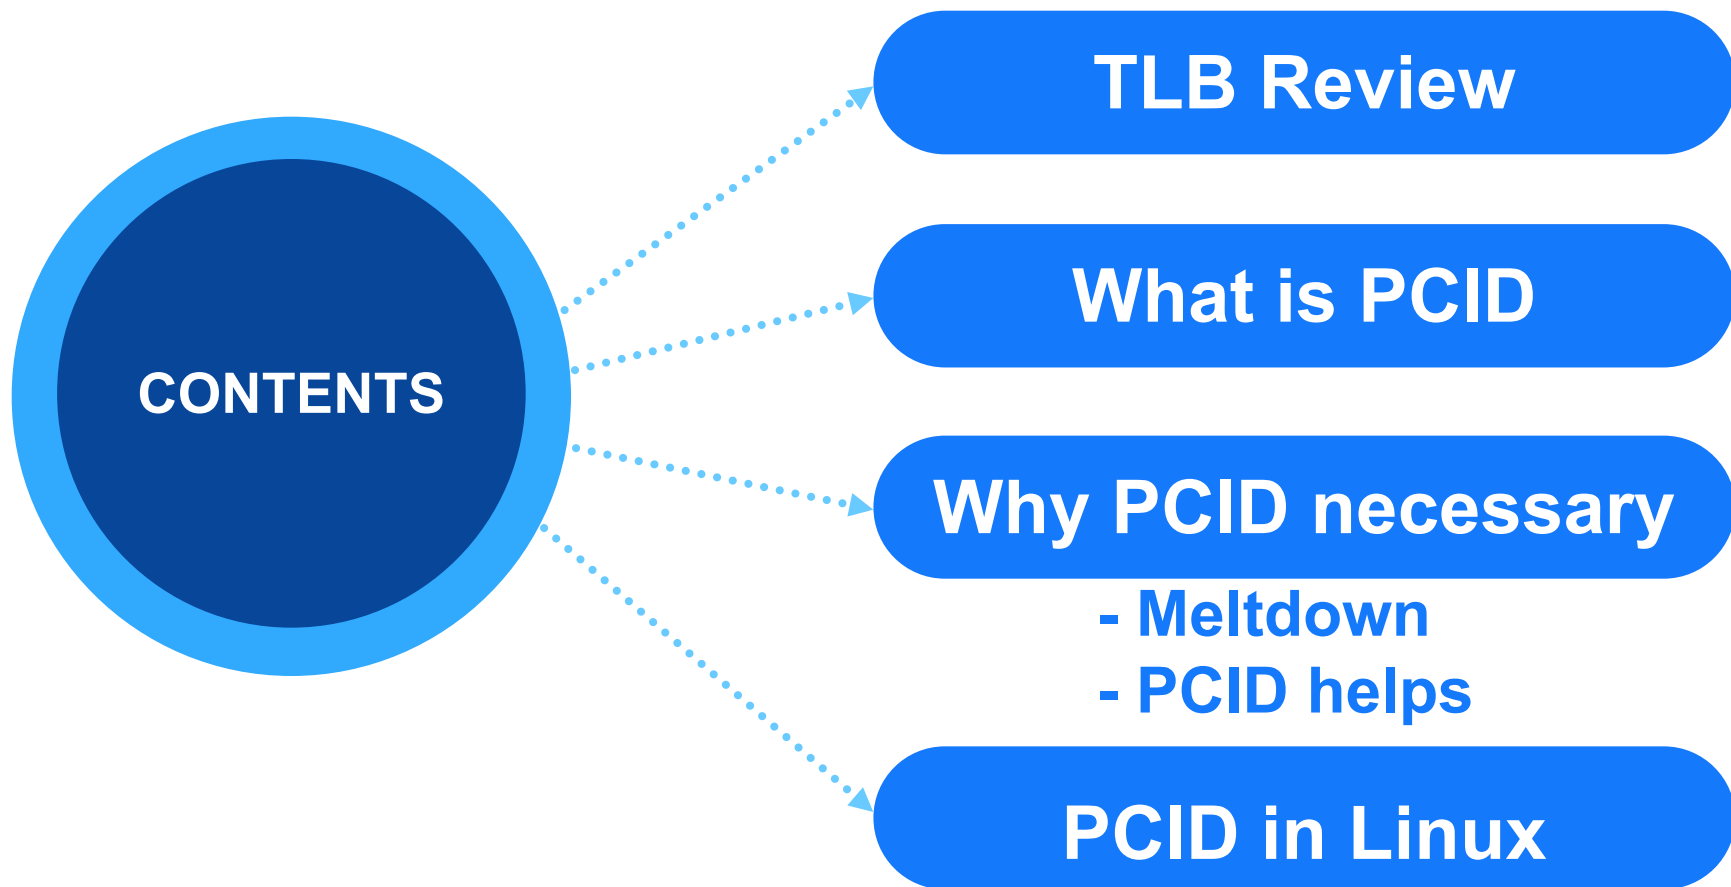


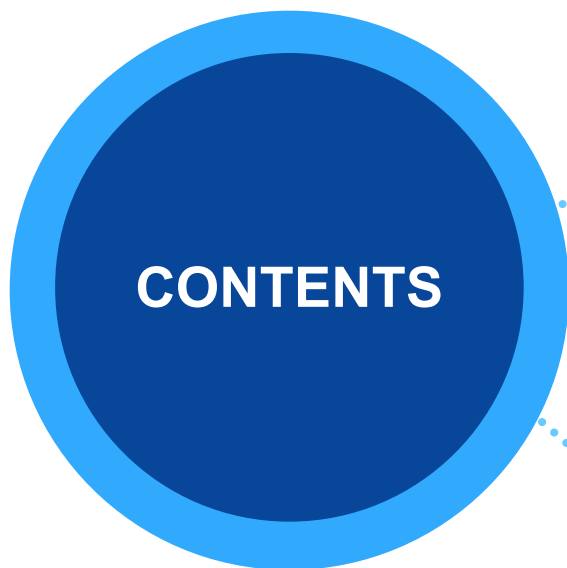
For CLK 2018

● PCID

a way to reduce task-switch cost

Zhaolei from RunCloud





CONTENTS

TLB Review

What is PCID

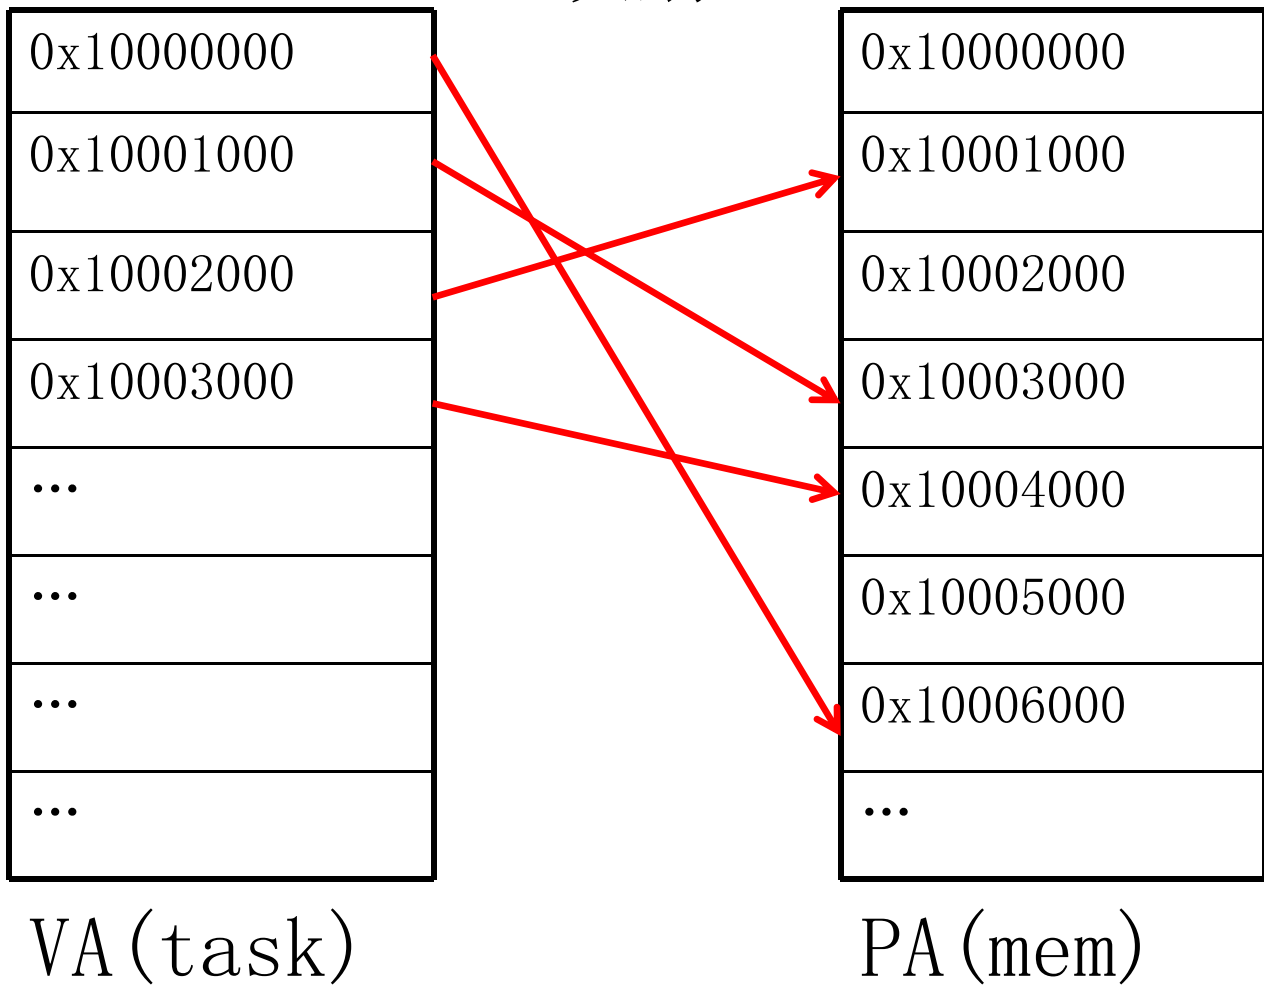
Why PCID necessary

- Meltdown
- PCID helps

PCID in Linux

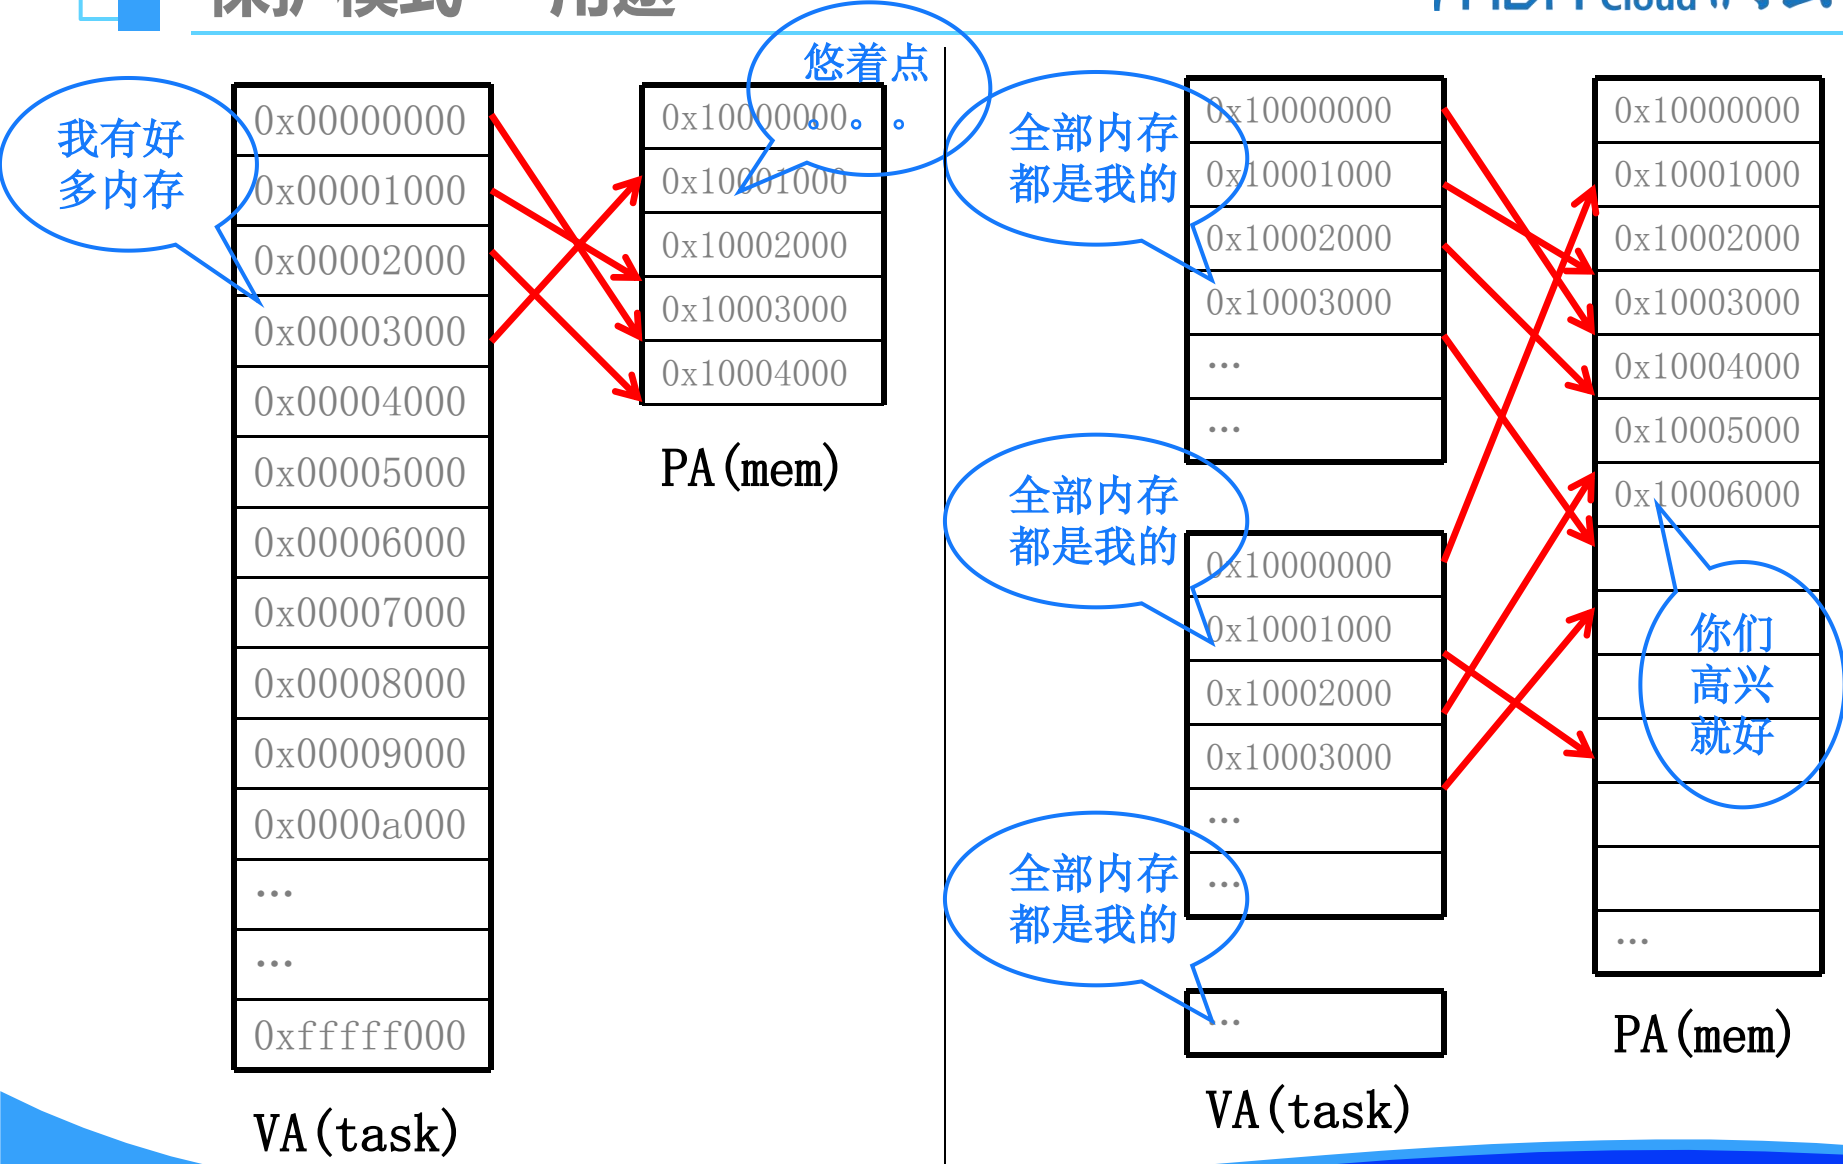


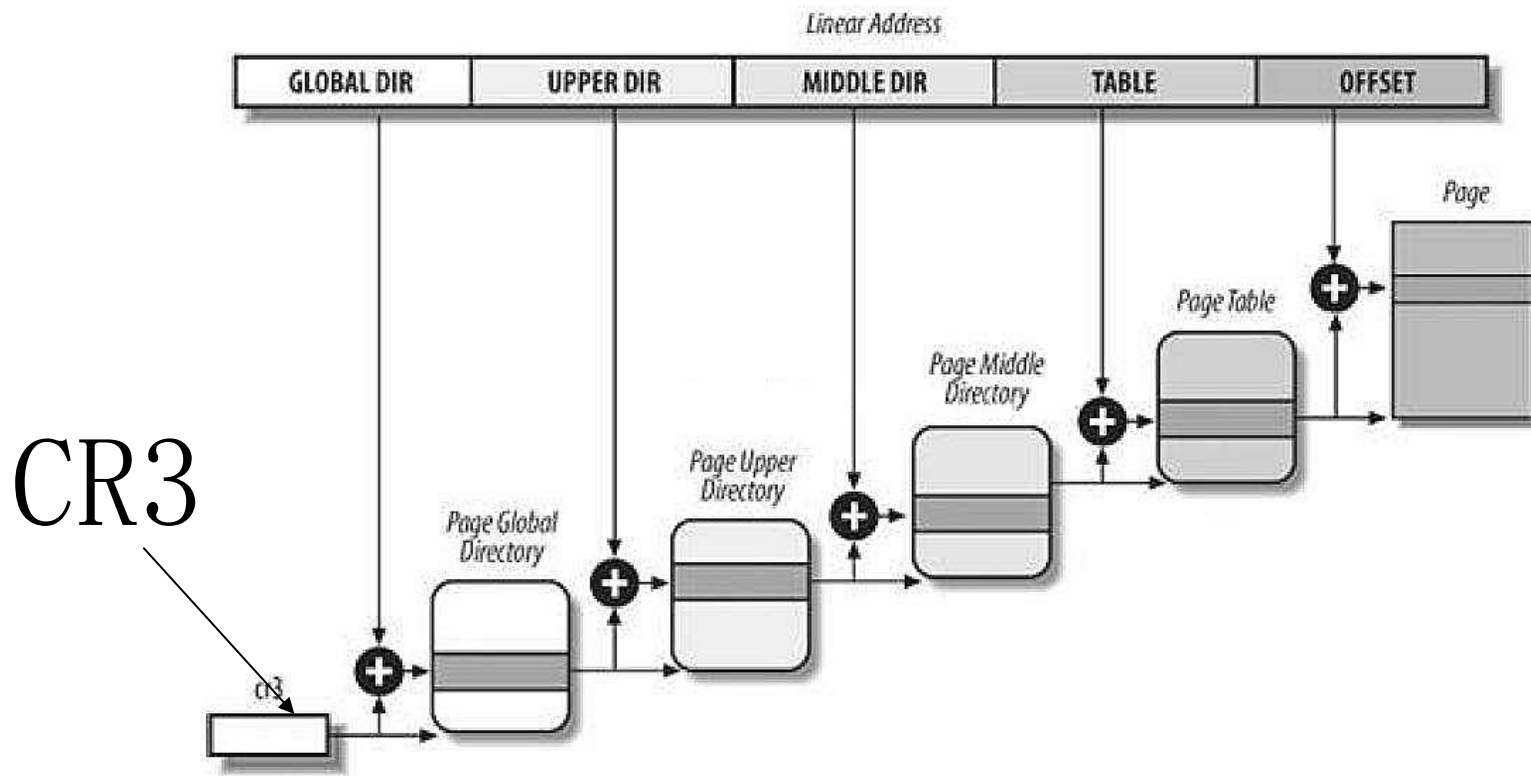
映射





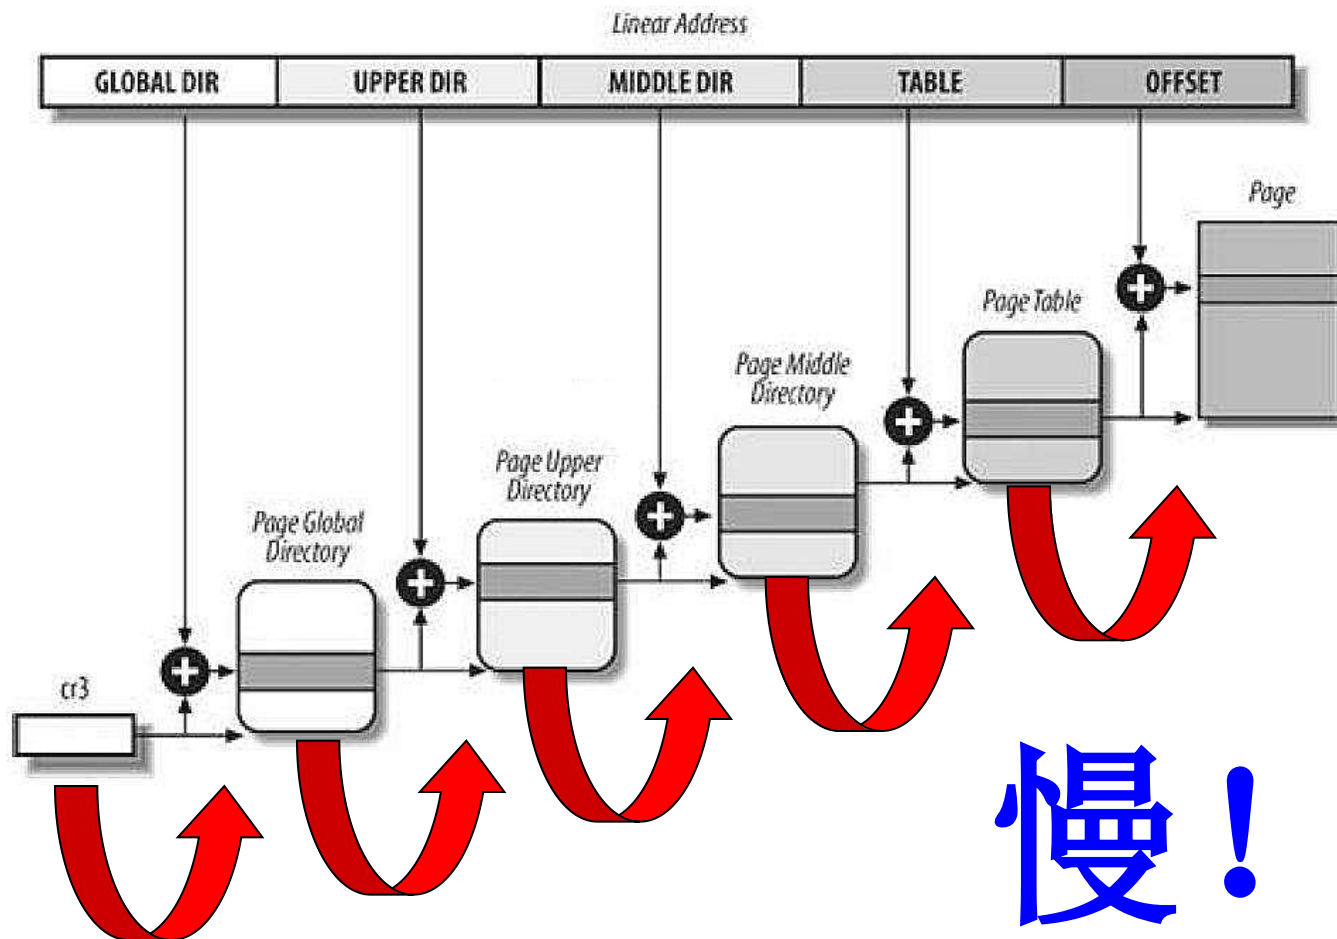
保护模式 - 用途





[mm_context.h](#)

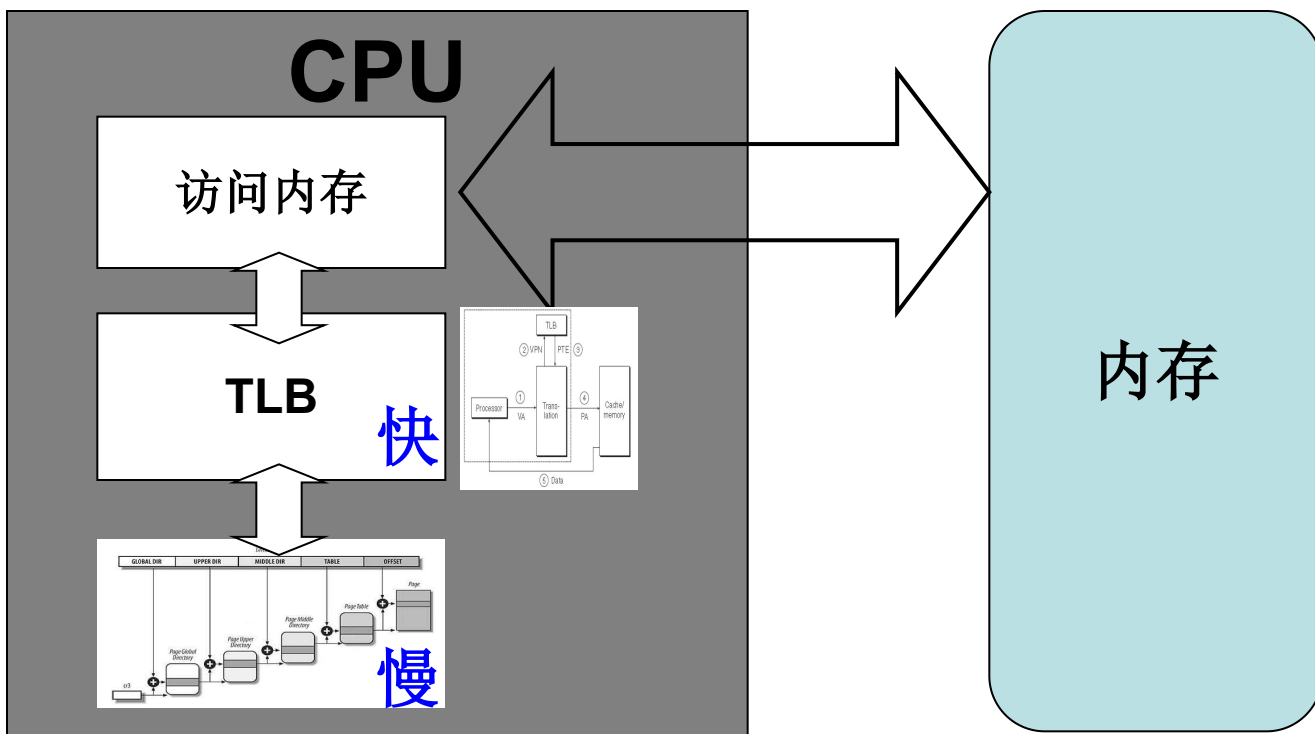
```
asm volatile("movl %0,%%cr3": : "r" (__pa(next->pgd));
```





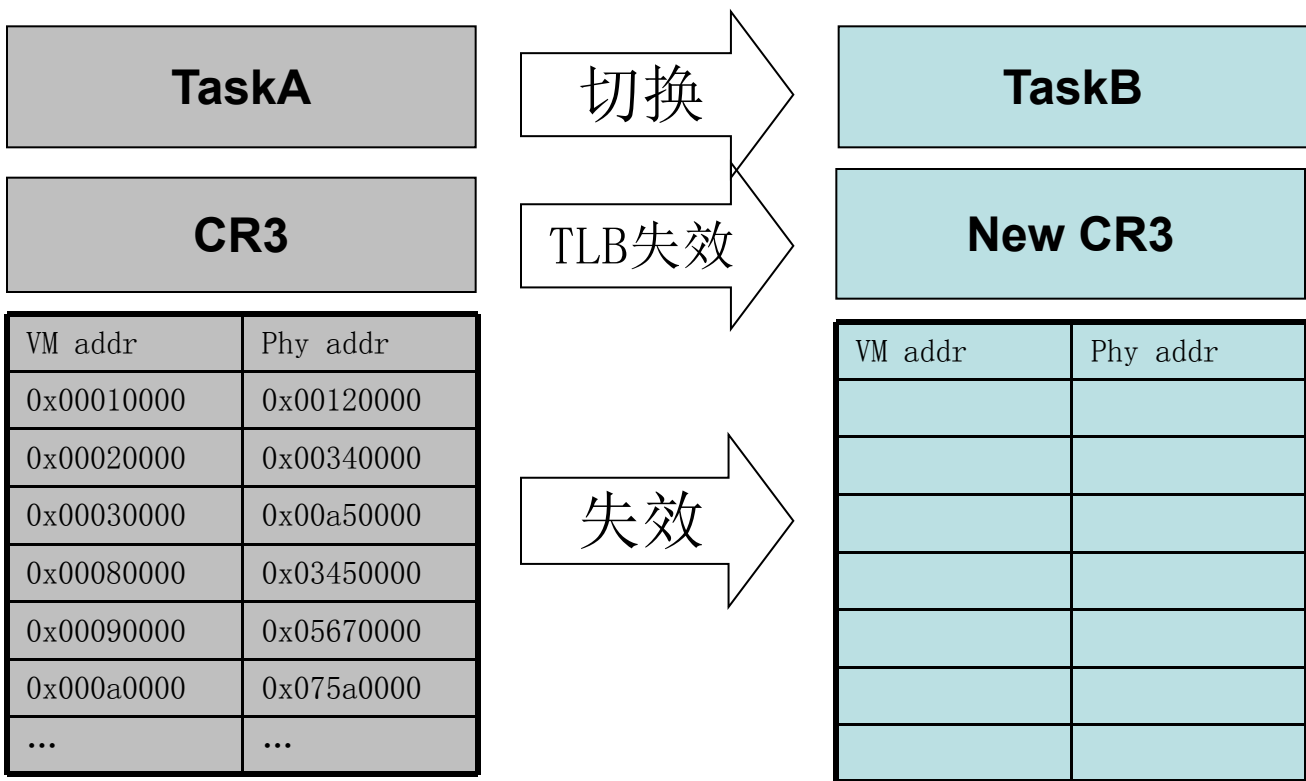
■ TLB

Translation Lookaside Buffer
页表缓冲、页表高速缓存





VM addr	Phy addr
0x00010000	0x00120000
0x00020000	0x00340000
0x00030000	0x00a50000
0x00080000	0x03450000
0x00090000	0x05670000
0x000a0000	0x075a0000
...	...



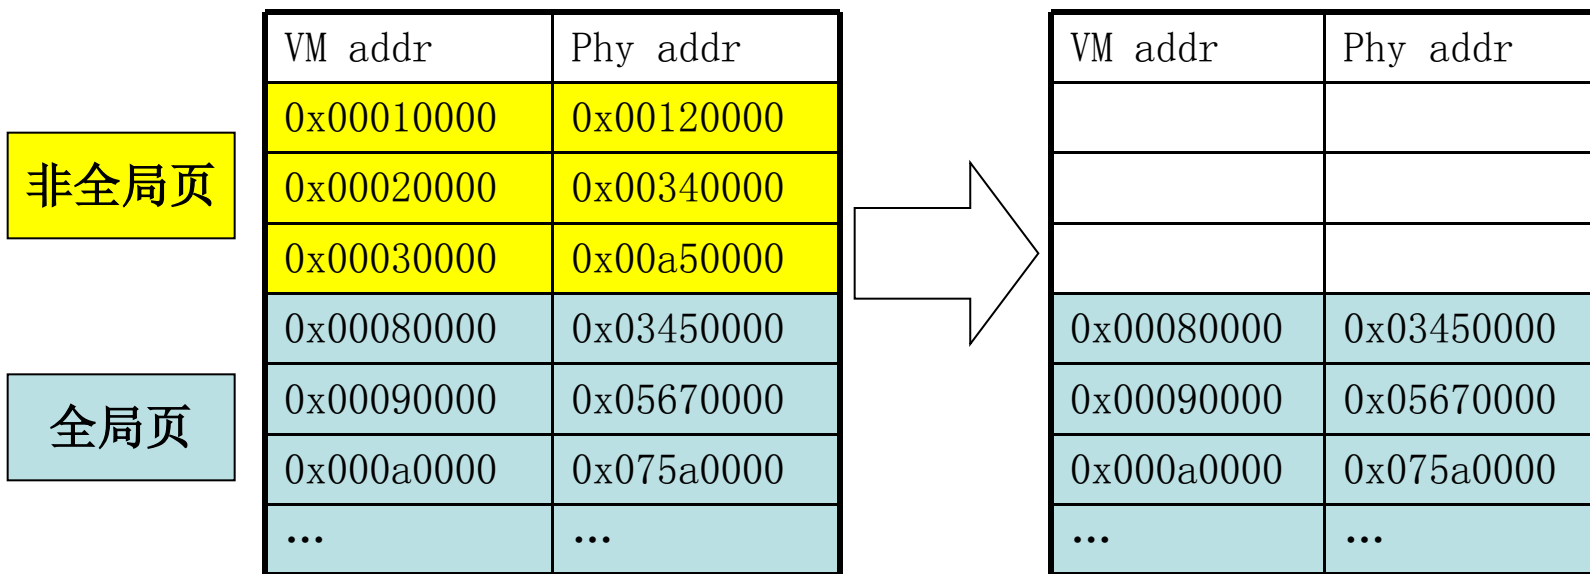
好不容易建立的缓存没了

`kernel/sched/core.c`

```
static inline struct rq * context_switch(struct rq *rq, struct task_struct *prev, struct task_struct *next)
```

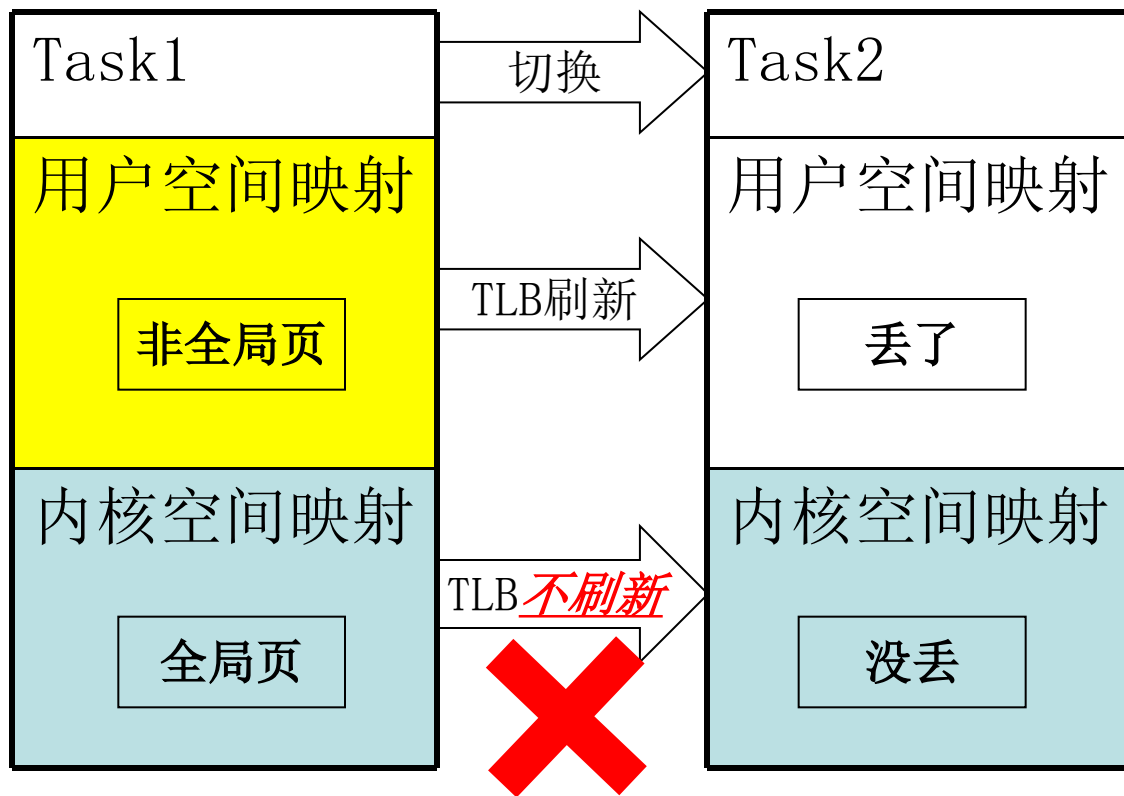


写入cr3 -> 失效非全局页的TLB项





写入cr3 -> 失效非全局页的TLB项

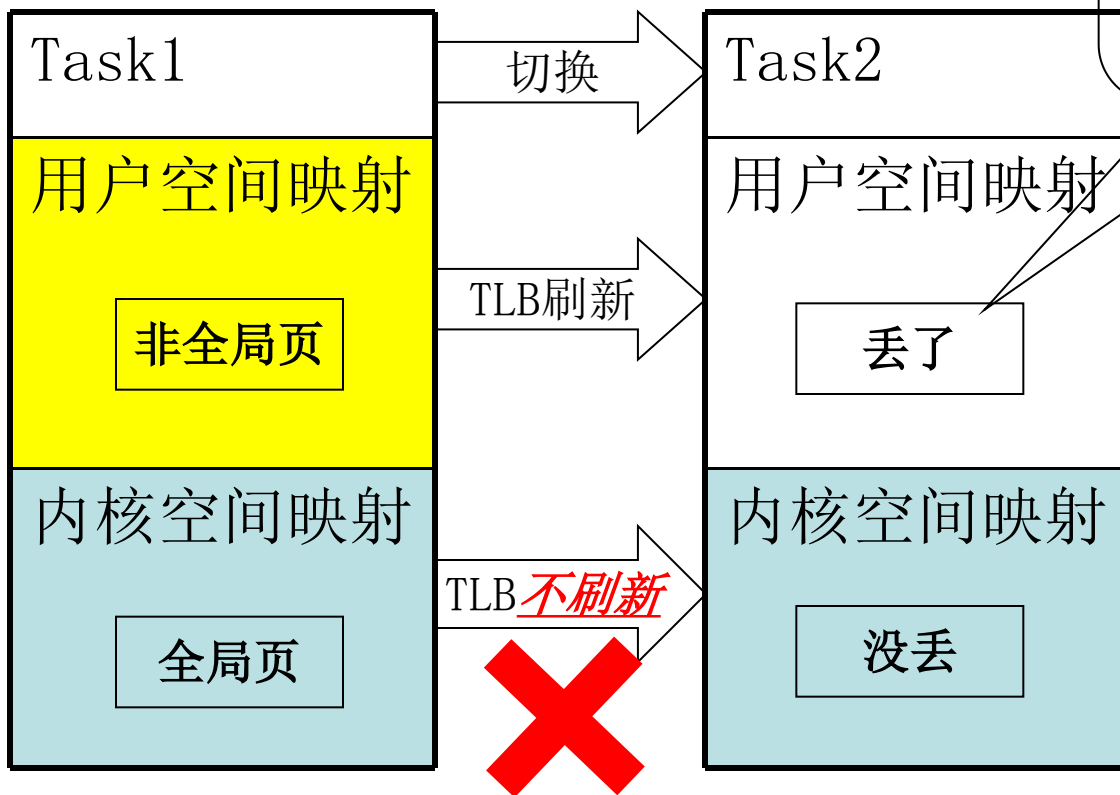


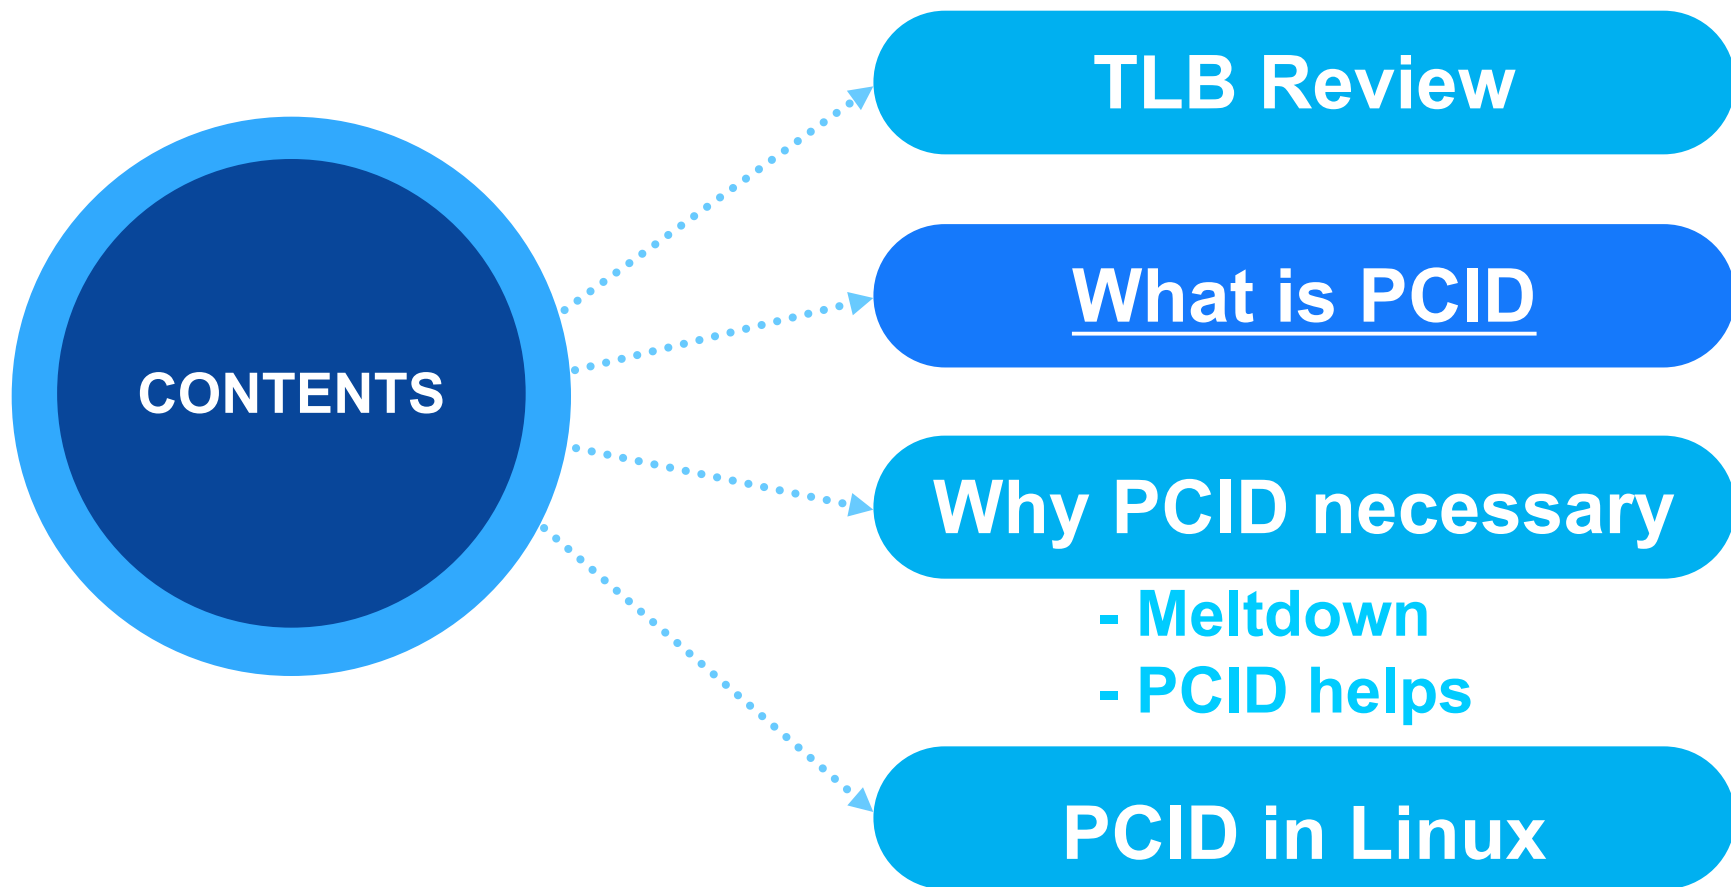


TLB的局限性 – 改善



这个也能不丢吗？







■ PCID - Processor Context ID

处理器上下文ID

■ From Westmere in 2010



这货还开始支持
1G大页了

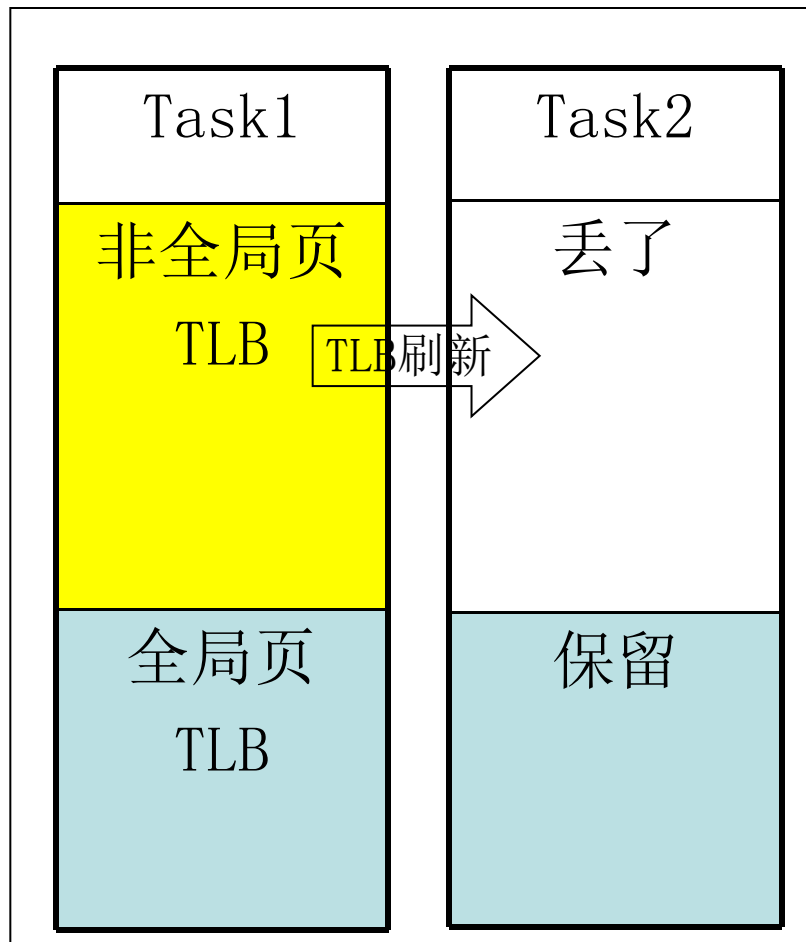


PCID

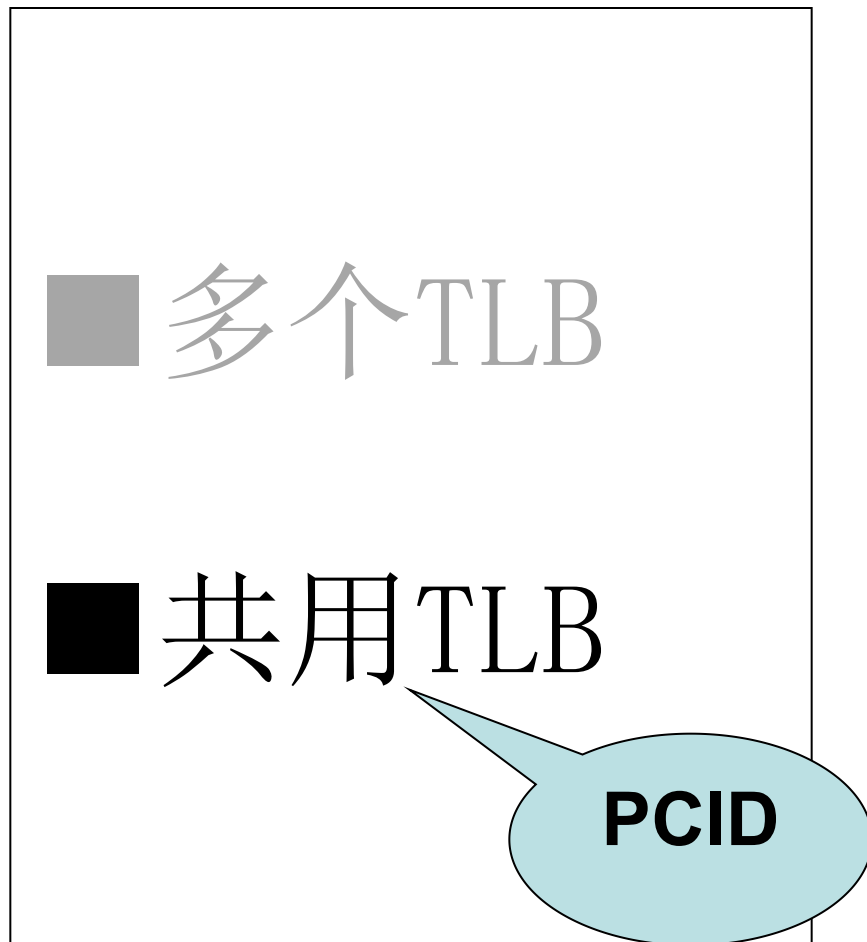
一句话用途
避免页表切换时的TLB丢失



Problem



Solution

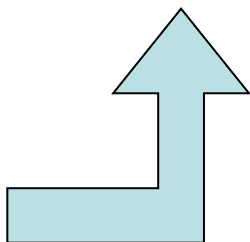




TLB with PCID

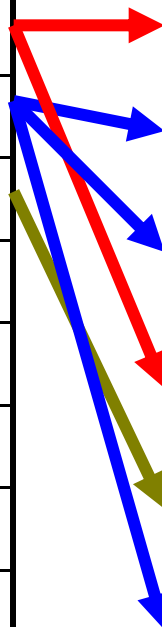
Index	VM addr	Phy addr
1	0x00010000	0x00120000
2	0x00020000	0x00340000
2	0x00030000	0x00a50000
1	0x00080000	0x03450000
3	0x00090000	0x05670000
2	0x000a0000	0x075a0000
...

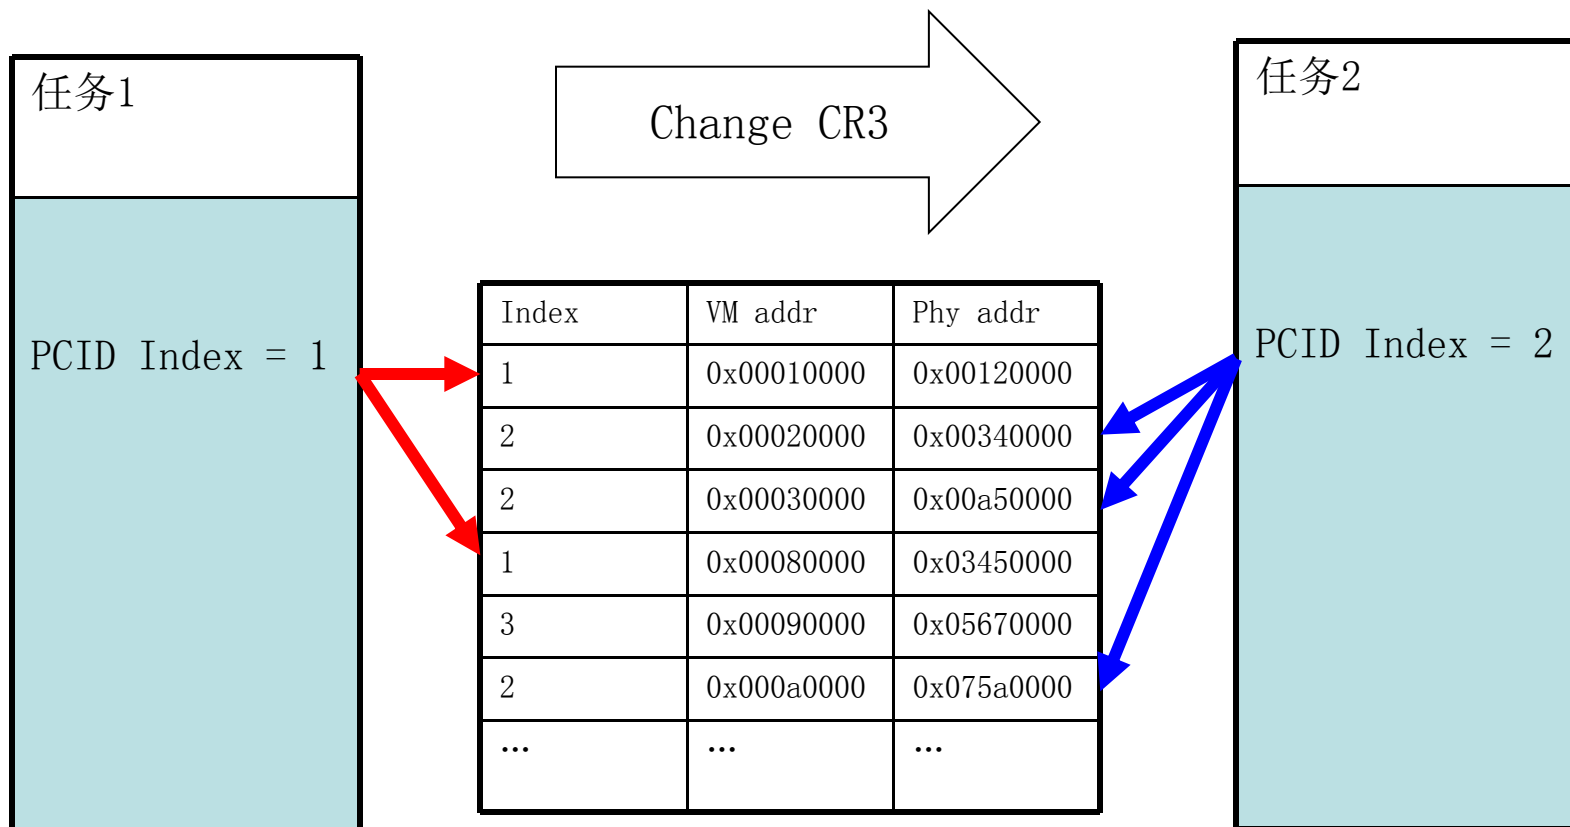
VM addr	Phy addr
0x00010000	0x00120000
0x00020000	0x00340000
0x00030000	0x00a50000
0x00080000	0x03450000
0x00090000	0x05670000
0x000a0000	0x075a0000
...	...



Task PCID Index
Index1
Index2
Index3
...
...
...
...
...
...

Index	VM addr	Phy addr
1	0x00010000	0x00120000
2	0x00020000	0x00340000
2	0x00030000	0x00a50000
1	0x00080000	0x03450000
3	0x00090000	0x05670000
2	0x000a0000	0x075a0000
...





CR4. PCIDE = 1

Index = CR3[0:11]

12bits, Max 4096 Items



PCID Introduced: 2010

PCID Supported by Linux: 2017

■ Reason1

4096 items limit, too many tasks

■ Reason2

Performance regression

■ Reason3

Not so many pagetable switch

https://kernelnewbies.org/Linux_4.14 (Longer-lived TLB Entries with PCID)

<https://zhuanlan.zhihu.com/p/32718446> (TLB shutdown)



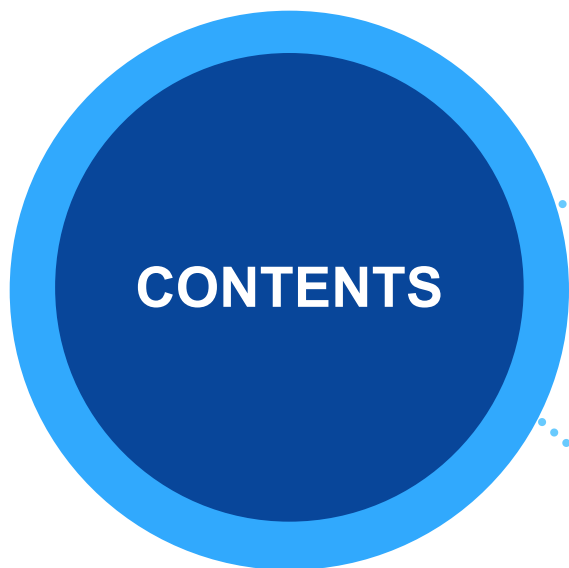
2018-1

What happened?

not very necessary

2018-1

~~not~~ very necessary



TLB Review

What is PCID

Why PCID necessary

- Meltdown
- PCID helps

PCID in Linux



What is Meltdown

2018-1-3

Google Project Zero (GPZ)

Jann Horn

CVE-2017-5715 – Spectre

CVE-2017-5753 – Spectre

CVE-2017-5754 – Meltdown



Meltdown



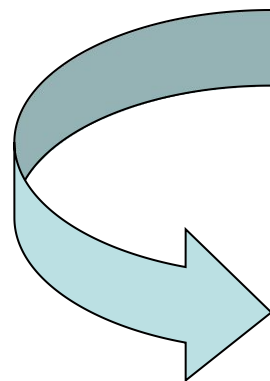
Spectre





代码:

```
a = 1;  
b = 2;  
c = 3;  
d = a + b + c;
```



乱序执行


缓存
侧信道

Since 1995
in Pentium Pro



执行:

```
a = 1;  
b = 2;  
c = 3;
```



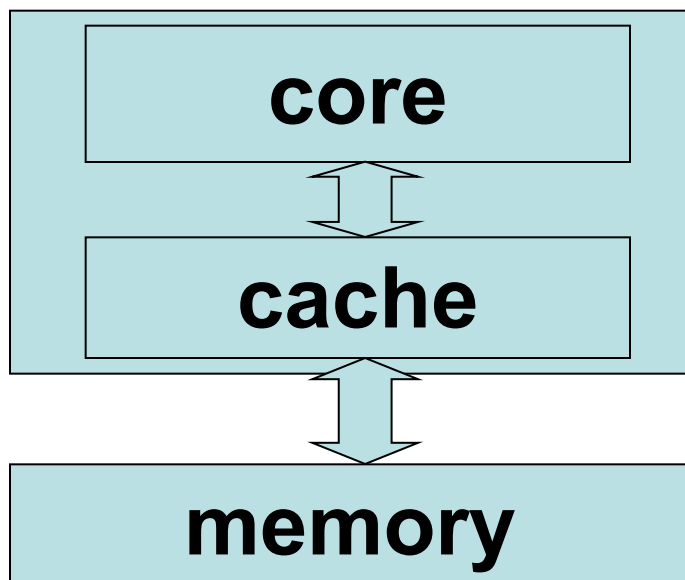
```
d = a + b + c;
```

流水线与乱序执行参见:

<https://blog.csdn.net/hyhop150/article/details/51440308>



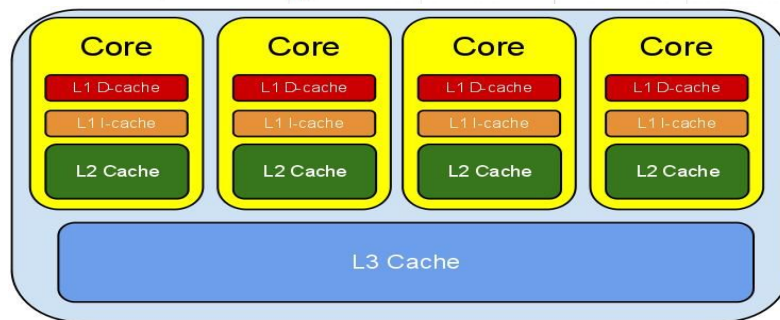
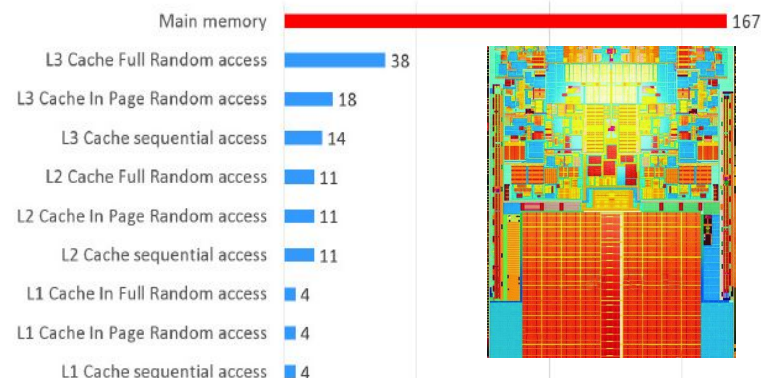
Meltdown原理



Processor Number	Cache	Clock Speed	Max TDP	Memory Type	Intel® HD Graphics	Number of Cores
i7-840QM	8 MB SmartCache	1.86 GHz	45 W	DDR3-1066/1333 MHz		4
i7-820QM	8 MB SmartCache	1.73 GHz	45 W	DDR3-1066/1333 MHz		4
i7-740QM	6 MB SmartCache	1.73 GHz	45 W	DDR3-1066/1333 MHz		4
i7-720QM	6 MB SmartCache	1.6 GHz	45 W	DDR3-1066/1333 MHz		4
i7-680UM	4 MB SmartCache	1.46 GHz	18 W	DDR3-800 MHz	✓	2
i7-660UM	4 MB SmartCache	1.33 GHz	18 W	DDR3-800 MHz	✓	2
i7-660UE	4 MB	1.33 GHz	18 W	DDR3-800 MHz	✓	2
i7-660LM	4 MB SmartCache	2.26 GHz	25 W	DDR3-800/1066 MHz	✓	2
i7-640UM	4 MB SmartCache	1.2 GHz	18 W	DDR3-800 MHz	✓	2
i7-640M	4 MB SmartCache	2.8 GHz	35 W	DDR3-800/1066 MHz	✓	2
i7-640LM	4 MB SmartCache	2.13 GHz	25 W	DDR3-800/1066 MHz	✓	2
i7-620UM	4 MB SmartCache	1.06 GHz	18 W	DDR3-800 MHz	✓	2
i7-620UE	4 MB SmartCache	1.06 GHz	18 W	DDR3-800 MHz	✓	2

乱序执行
缓存
侧信道

CPU Cache Access Latencies in Clock Cycles





Meltdown原理

```
int i;  
char array[256][4096];
```

乱序执行
缓存
侧信道

```
-----  
_mm_clflush(&target_array, sizeof(array));  
-----
```

清空测试数组的
CPU 缓存

```
// Normal instructions
```

```
array[*ADDR_TO_READ][0] = 1;
```

向测试数组写入
探测数据

```
-----  
// Trap to segfault
```

```
for (i = 0; i < 256; i++) {  
    time_begin = rdtsc();  
    j = array[i][0];  
    time_end = rdtsc;  
    if (time_end - time_start < THRESHOLD)  
        print("Data is %d\n", i);  
}
```

根据探测数据的
访问时间，得到
希望的数据



```
1) // Normal instructions  
2) array[*ADDR_TO_READ][0] = 1;
```

乱序执行
缓存
侧信道

一般理解的运行过程:

```
exec(1) -> check_perm(2) -> failed -> segfault
```

CPU的实际运行过程:

```
exec(1)  
    -> check_perm(2) -> failed -> rollback(2) -> segfault  
exec(2)
```

结果: 数据没变, 缓存变了

array[ADDR_TO_READ][0] in cache



```
for (i = 0; i < 256; i++) {  
    time_begin = rdtsc();  
    j = array[i][0];  
    time_end = rdtsc();  
    if (time_end - time_start < THRESHOLD)  
        print("Data is %d\n", i);  
}
```

乱序执行
缓存
侧信道

```
array[0][0] - 20 cycles  
array[1][0] - 24 cycles  
array[2][0] - 19 cycles  
array[3][0] - 23 cycles  
array[4][0] - 25 cycles  
...  
array[ADDR_TO_READ][0] - 2 cycles  
array[1][0] - 22 cycles  
array[1][0] - 19 cycles
```

之前的操作让**array[ADDR_TO_READ][0]**进入了缓存，
所以这个数据读取的会比其他数据稍微快一些

采用高精度计时器(**TSC**等)衡量操作的时间，即可知道
ADDR_TO_READ的内容



乱序执行
缓存
侧信道

Meltdown



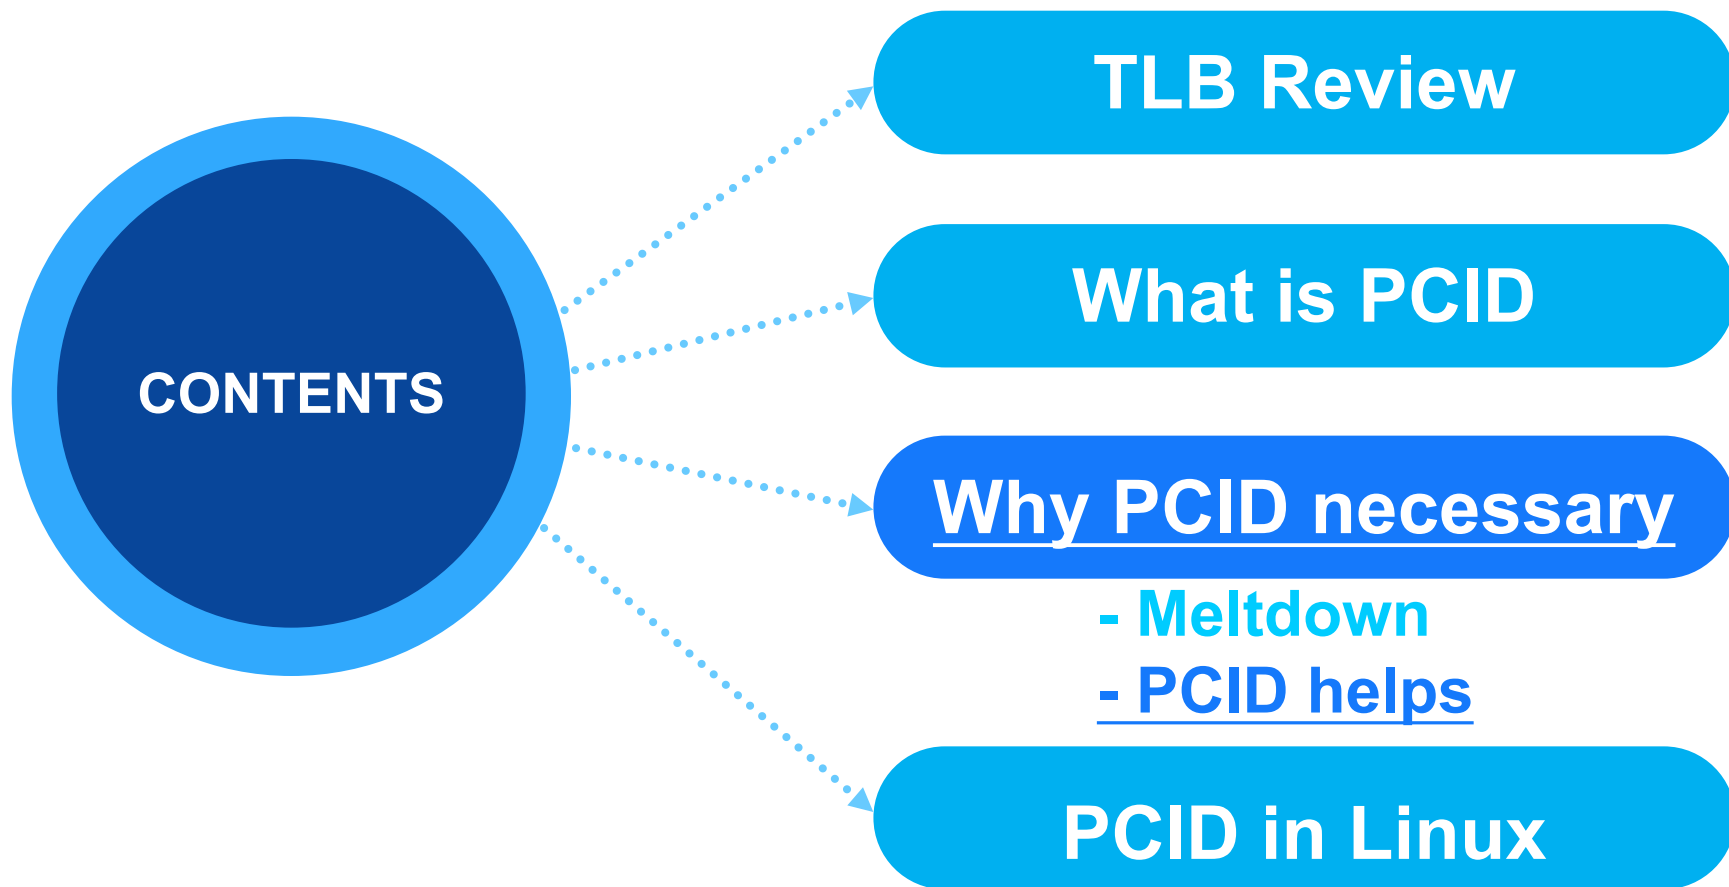
分支预测
缓存
侧信道

Spectre

参见:

一步一步理解CPU芯片漏洞: Meltdown与Spectre

<http://www.freebuf.com/articles/system/159811.html>





meltdown-exploit

<https://github.com/paboldin/meltdown-exploit>

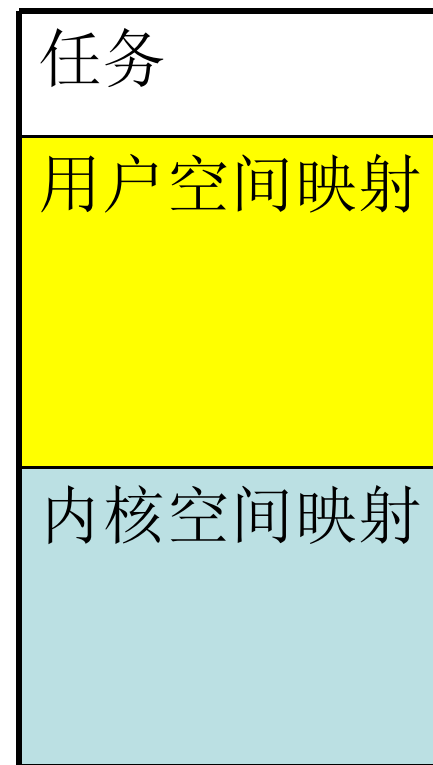
Commit	Author	Message	Date
47 commits			
Latest commit	36e17e7	Feb 13, 2018	
Branch: master	New pull request	Create new file	Upload files
Find file	Clone or download		
gitignore	detect rotscp instruction at completime	Jan 17, 2018	
Makefile	detect rotscp instruction at completime	Jan 17, 2018	
README.md	reference full exploit	Jan 19, 2018	
detect_rotscp.sh	detect rotscp instruction at completime	Jan 17, 2018	
meltdown.c	Add mfence instruction before speculation	Feb 14, 2018	
run.sh	read /lib/modules/\${uname -r}/build/System.map	Jan 31, 2018	
README.md			

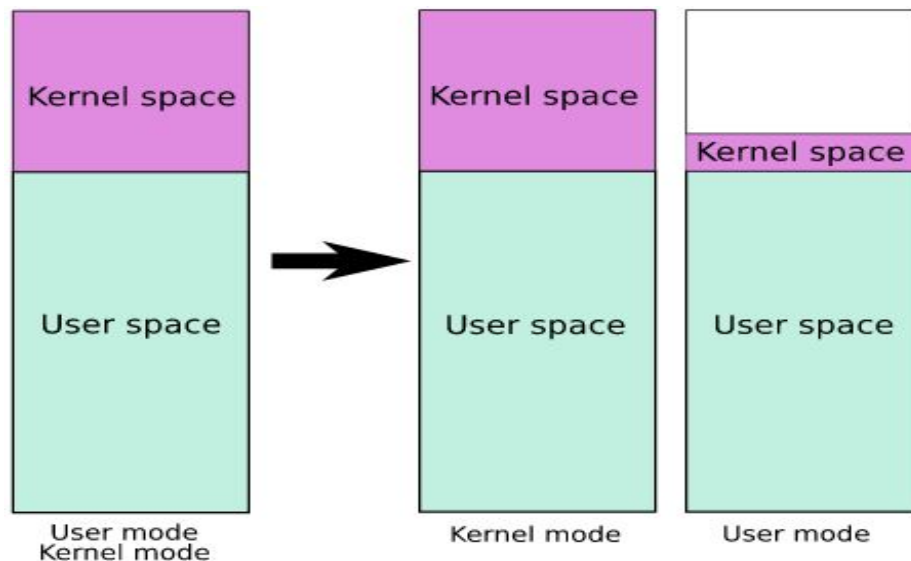
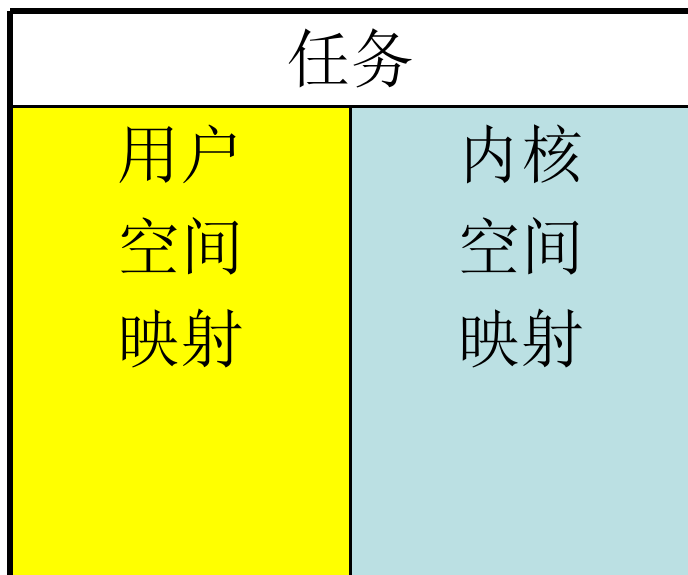
```
$ ./run.sh
looking for linux_proc_banner in /proc/kallsyms
protected. requires root
+ find_linux_proc_banner /proc/kallsyms sudo
+ sudo awk

        /linux_proc_banner/ {
            if (strtonum("0x"$1))
                print $1;
            exit 0;
        } /proc/kallsyms
+ linux_proc_banner=ffffffffffa3e000a0
+ set +x
cached = 29, uncached = 271, threshold 88
read ffffffff3e000a0 = 25 %
read ffffffff3e000a1 = 73 s
read ffffffff3e000a2 = 20
read ffffffff3e000a3 = 76 v
read ffffffff3e000a4 = 65 e
read ffffffff3e000a5 = 72 r
read ffffffff3e000a6 = 73 s
read ffffffff3e000a7 = 69 i
read ffffffff3e000a8 = 6f o
read ffffffff3e000a9 = 6e n
read ffffffff3e000aa = 20
read ffffffff3e000ab = 25 %
read ffffffff3e000ac = 73 s
read ffffffff3e000ad = 20
read ffffffff3e000ae = 28 (
read ffffffff3e000af = 62 b
read ffffffff3e000b0 = 75 u
read ffffffff3e000b1 = 69 i
read ffffffff3e000b2 = 6c l
read ffffffff3e000b3 = 64 d
read ffffffff3e000b4 = 64 d
read ffffffff3e000b5 = 40 @
VULNERABLE
VULNERABLE ON
4.10.0-42-generic #
```



- 目标数据需要与探测代码处于同一地址空间
- 传统内核因为性能考虑，把内核空间与用户空间映射在同一地址空间





KAISER

kernel address space layout randomization

KPTI

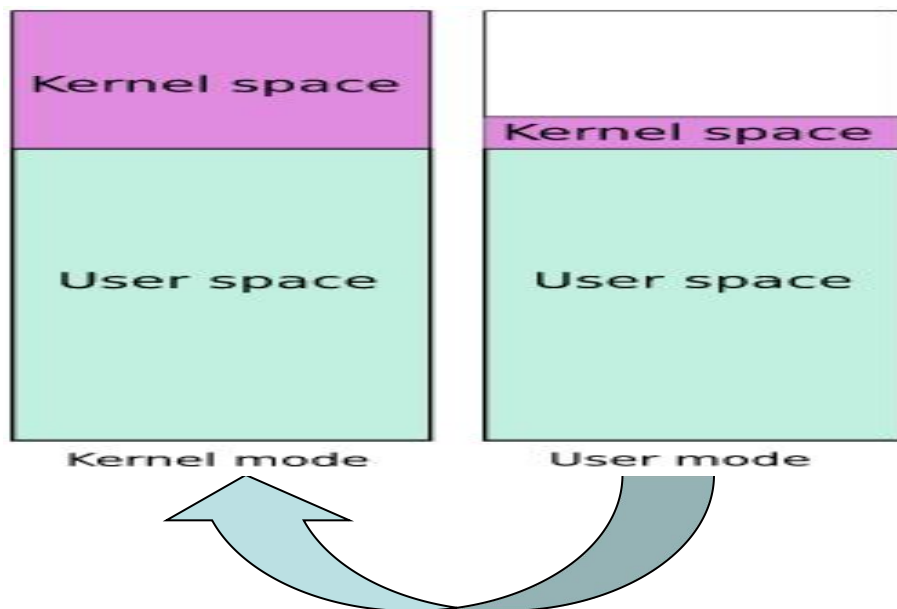
Kernel page-table isolation

arch/x86/mm/pti.c

```
/*  
 * Initialize kernel page table isolation  
 */  
void __init pti_init(void)  
{  
    if (!static_cpu_has(X86_FEATURE_PTI))  
        return;  
  
    pr_info("enabled\n");  
  
#ifdef CONFIG_X86_32  
    /*  
     * We check for X86_FEATURE_PCID here. But the init-code will  
     * clear the feature flag on 32 bit because the feature is not  
     * supported on 32 bit anyway. To print the warning we need to  
     * check with cpuid  
     */  
#endif  
}
```



Meltdown防御的性能问题



Pagetable switch cnt:

BEFORE: = task_switch_cnt

AFTER: = task_switch_cnt
+ syscall_cnt
+ interrupt_cnt
= N * BEFORE

```
[root@RUNCLOUD_ZL]# wc -l /tmp/context_switch
[root@RUNCLOUD_ZL]# 4743 /tmp/context_switch
[root@RUNCLOUD_ZL]#
[root@RUNCLOUD_ZL]# wc -l /tmp/sys_enter
[root@RUNCLOUD_ZL]# 1364787 /tmp/sys_enter
[root@RUNCLOUD_ZL]#
```

1364787 / 4743 = 287

* Test result of ftrace in a host(5s)



■ Problem

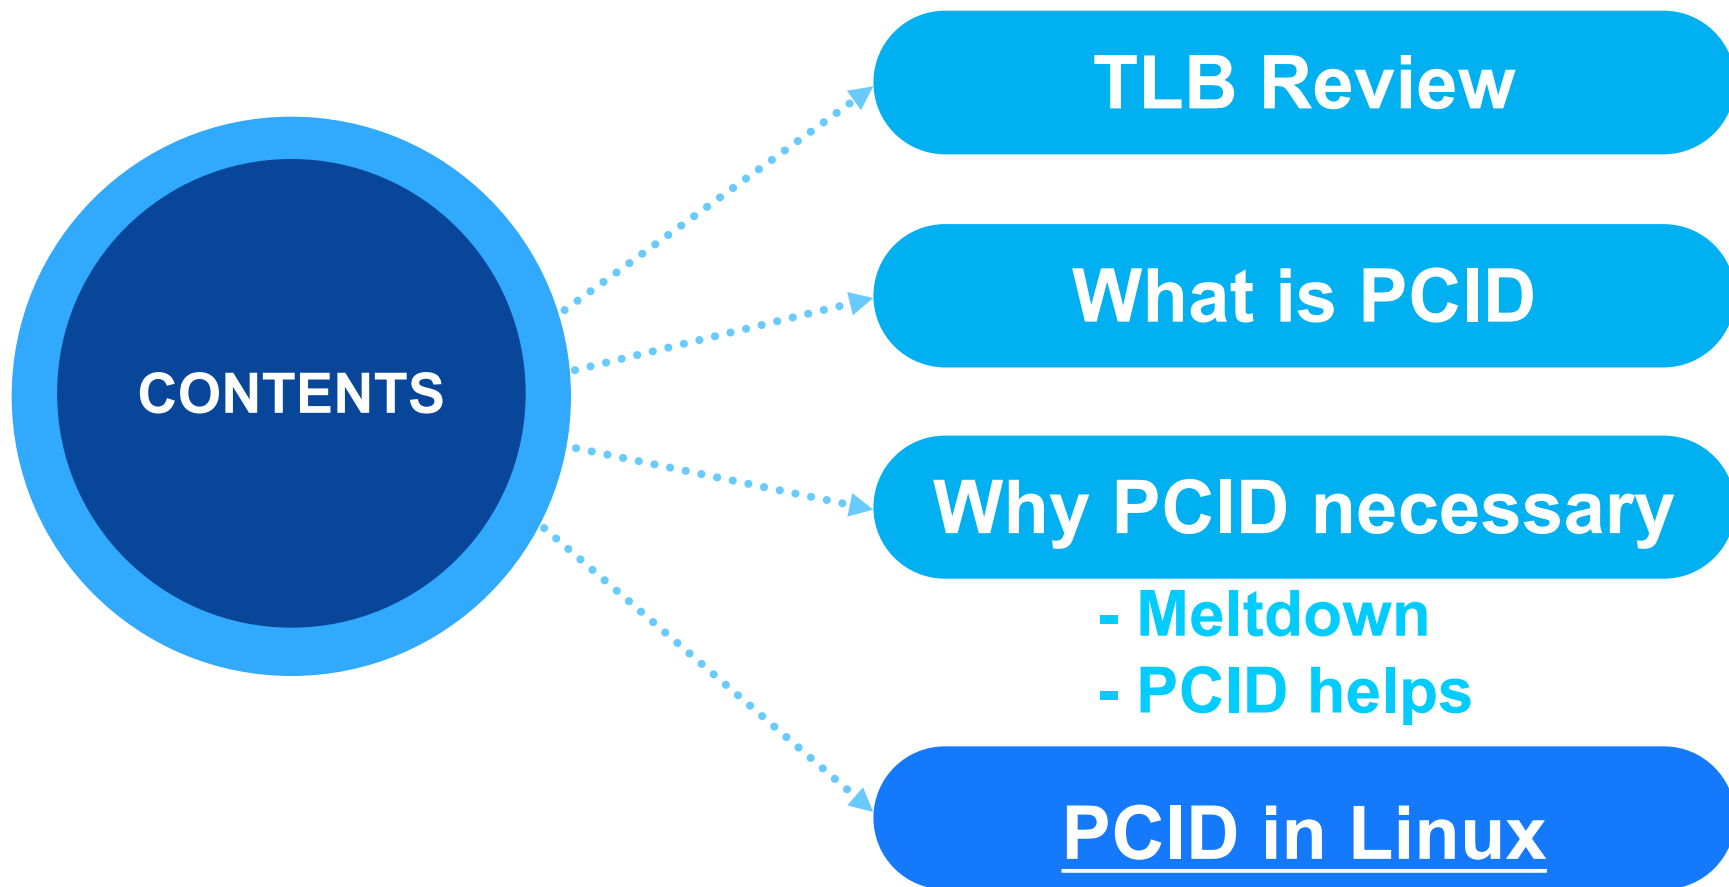
pagetable switch * N

→ Flush TLB * N

→ Too many cache lost

■ Solution

Use PCID to reduce cache lost
caused by pagetable switch.





commit d3b5d35290d729a2518af00feca867385a1b08fa
Merge: aa2a4b65 7138970
Author: Linus Torvalds <torvalds@linux-foundation.org>
Date: **Mon May 1 23:54:56 2017 -0700**

Merge branch 'x86-mm-for-linus' of git://git.kernel.org/pub/scm/linux/kernel/git/tip/tip

V4.12

Pull x86 mm updates from Ingo Molnar:

"The main x86 MM changes in this cycle were:

- **continued native kernel PCID support preparation patches to the TLB flushing code (Andy Lutomirski)**

commit 7a69f9c60b49699579f5bfb71f928cceb0afe1a
Merge: 9bc088a 8781fb7
Author: Linus Torvalds <torvalds@linux-foundation.org>
Date: **Mon Jul 3 14:45:09 2017 -0700**

Merge branch 'x86-mm-for-linus' of git://git.kernel.org/pub/scm/linux/kernel/git/tip/tip

Pull x86 mm updates from Ingo Molnar:

"The main changes in this cycle were:

V4.13

...

- **Continued work to add PCID CPU support to native kernels as well. In this round most of the focus is on reworking/refreshing the TLB flush infrastructure for the upcoming PCID changes. (Andy Lutomirski)"**



commit b1b6f83ac938d176742c85757960dec2cf10e468

Merge: 5f82e71 9e52fc2

Author: Linus Torvalds <torvalds@linux-foundation.org>

Date: **Mon Sep 4 12:21:28 2017 -0700**

Merge branch 'x86-mm-for-linus' of git://git.kernel.org/pub/scm/linux/kernel/git/tip/tip

v4.14

Pull x86 mm changes from Ingo Molnar:

"PCID support, 5-level paging support, Secure Memory Encryption support

...

- **Enable PCID optimized TLB flushing on newer Intel CPUs: PCID is a hardware feature that attaches an address space tag to TLB entries and thus allows to skip TLB flushing in many cases, even if we switch mm's.**

(By Andy Lutomirski)



```
[root@RUNCLOUD_ZL linux]# git log --oneline | grep -w -i pcid
```

88c6f8a x86/mm/pti: Fix 32 bit PCID check

5e81059 x86/mm/pti: Add Warning when booting on a PCID capable CPU **v4.19**

8c06c77 x86/pti: Leave kernel text global for !PCID **v4.17**

f10ee3d x86/pti: Fix !PCID and sanitize defines

0a126ab x86/mm: Clarify the whole ASID/kernel PCID/user PCID naming

6fd166a x86/mm: Use/Fix PCID to optimize user/kernel switches

fae1a3e kvm: x86: fix RSM when PCID is non-zero

52a2af4 x86/mm/64: Stop using CR3.PCID == 0 in ASID-aware code **v4.15**

f34902c x86/hibernate/64: Mask off CR3's PCID bits in the saved CR3 **v4.14**

7898f79 x86/mm/64: Fix an incorrect warning with CONFIG_DEBUG_VM=y, !PCID

10af623 x86/mm: Implement PCID based optimization: try to preserve old TLB entries using PCID

0790c9a x86/mm: Add the 'nopcid' boot option to turn off PCID

cba4671 x86/mm: Disable PCID on 32-bit kernels

```
[root@RUNCLOUD_ZL linux]#
```




```
[root@RUNCLOUD_ZL]# vi Documentation/admin-guide/kernel-parameters.txt
```

```
...
```

nopcid [X86-64] Disable the PCID cpu feature.

noinvpcid [X86] Disable the INVPCID cpu feature.

```
...
```

```
[root@RUNCLOUD_ZL]#
```

```
[root@RUNCLOUD_ZL]# vi x86/pti.txt
```

h. INVPCID is a TLB-flushing instruction which allows flushing of TLB entries for non-current PCIDs. Some systems support PCIDs, but do not support INVPCID. On these systems, addresses can only be flushed from the TLB for the current PCID. When flushing a kernel address, we need to flush all PCIDs, so a single kernel address flush will require a TLB-flushing CR3 write upon the next use of every PCID.

```
[root@RUNCLOUD_ZL]#
```



Thanks



Thanks