



# How to add SMP for a New Arch

Guo Ren

<https://github.com/c-sky/csky-linux>

## abiv1

- **16b ISA**, 32b data
- [M-Core compatible](#)
- **Mass production**

### Product:

- 510
- 610 (linux supported)

## abiv2

- EM\_CSKY 252
- **32b/16b mixed ISA**, 32b data
- 1st independent ISA/Arch
- **Mass production**

### Product:

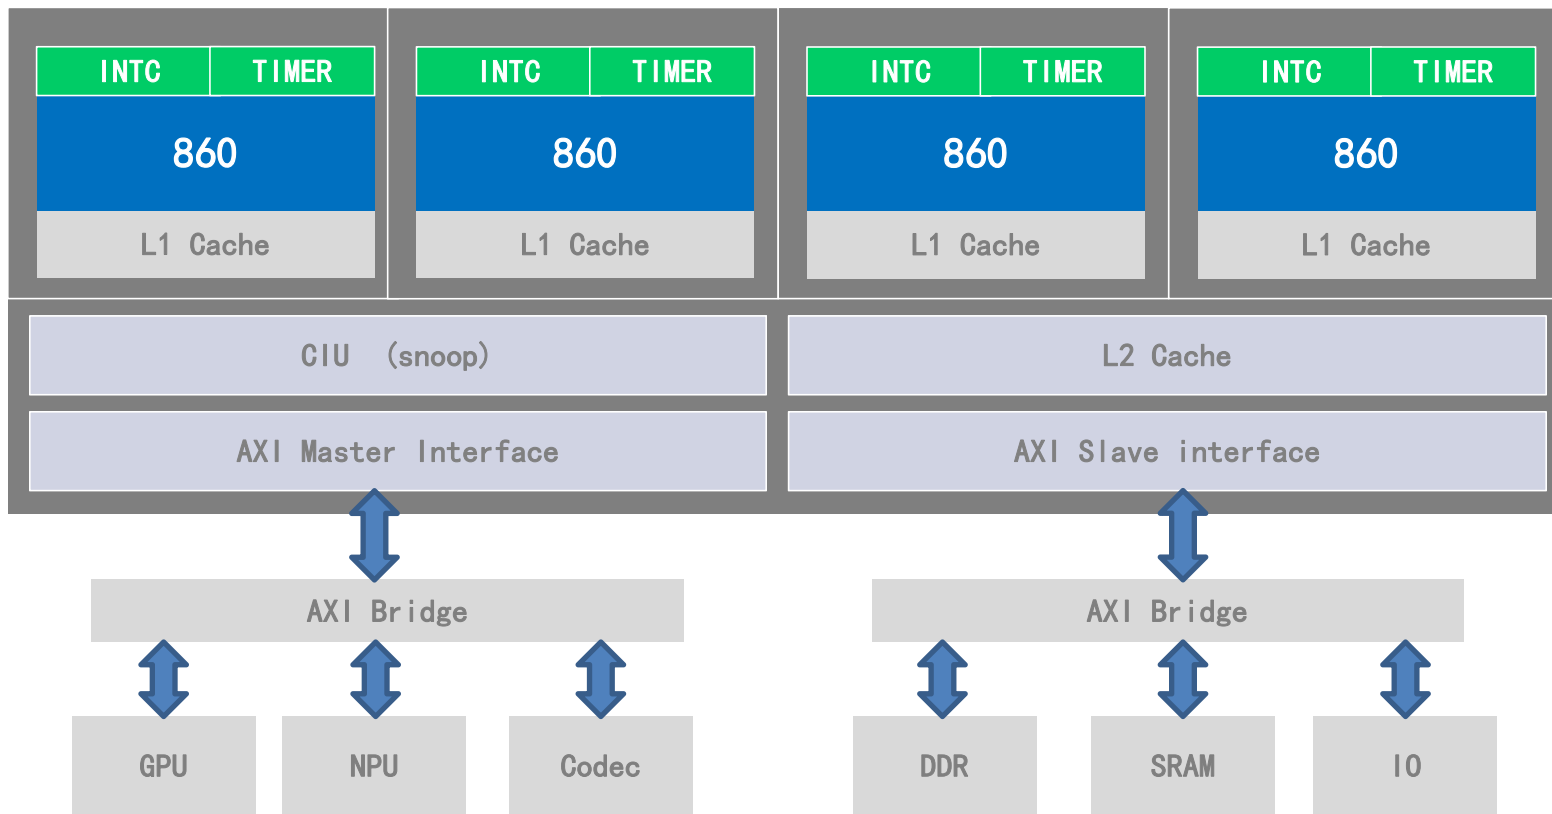
- **807 (7-pipeline linux supported)**
- **810 (10-pipeline linux supported)**
- **860 (12-pipeline linux supported)**
- **F: FPU (V3 ISA)**
- **V: VDSP (V3 ISA SIMD-128bit)**
- **T: Trust Exe Environment**
- **SMP supported**

-mcpu=

807f 807fv 807ft 807fvt

810f 810fv 810ft 810fvt

860f 860fv 860ft 860fvt



## abiv1

- **16b ISA**, 32b data
- [M-Core compatible](#)
- **Mass production**

### Product:

- 510
- 610 (linux supported)

## abiv2

- EM\_CSKY 252
- **32b/16b mixed ISA**, 32b data
- 1st independent ISA/Arch
- **Mass production**

### Product:

- **807 (7-pipeline linux supported)**
- **810 (10-pipeline linux supported)**
- **860 (12-pipeline linux supported)**
- **F: FPU (V3 ISA)**
- **V: VDSP (V3 ISA SIMD-128bit)**
- **T: Trust Exe Environment**
- **SMP supported**

### -mcpu=

807f 807fv 807ft 807fvt  
810f 810fv 810ft 810fvt  
860f 860fv 860ft 860fvt

## abiv3 ? (No yet born)

- Making love with RISC-V

## How to add SMP for a new arch

- SMP Boot
- MP-INTC Introduction (IPI mechanism)
- MP-TIMER Introduction
- Atomic implementation
- Spinlock, rwlock (test-and-set V.S ticket-queue)
- Memory Barrier
- Cache & TLB sync for SMP

## SMP Qemu demo

## Q & A

- soc-reset-control V.S. free-run
- C-SKY RMR (Reset Manage Register)

Why use reset-controller?

- make SMP boot simpler
- make ck807/ck810/ck860 use one bootroom/bootloader binary, one work for ever

Bootloader:

[https://github.com/c-sky/buildroot/blob/master/board/csky/ck860\\_platform/gdbinit](https://github.com/c-sky/buildroot/blob/master/board/csky/ck860_platform/gdbinit)

Linux head.S:

<https://github.com/c-sky/csky-linux/blob/master/arch/csky/kernel/head.S>

Linux csky\_start -> start\_kernel()

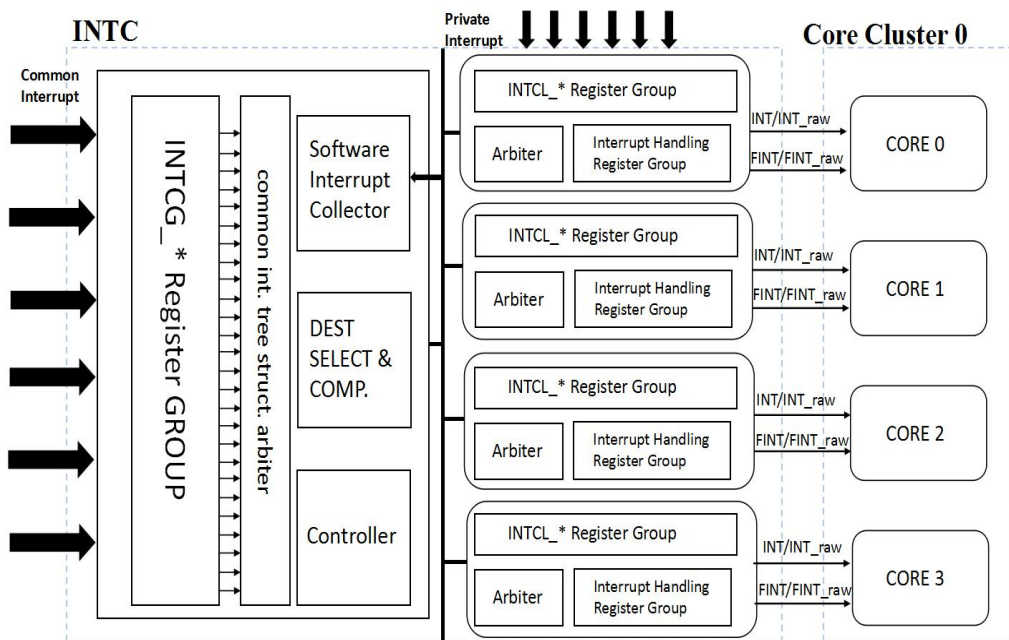
<https://github.com/c-sky/csky-linux/blob/master/arch/csky/kernel/setup.c>

- 0 - 15 : software irq, and we use 15 as our IPI\_IRQ.
- 16 - 31 : private irq, and we use 16 as the co-processor timer.
- 31 - 1024 : external irq for soc ip.

- Nice Software interface (window-reg access)
- 256 levels IRQ priority setting (8-bits width)
- Support 4 IRQ trigger types config:

1. LEVEL\_HIGH
2. LEVEL\_LOW
3. EDGE\_RISING
4. EDGE\_FALLING

- (Determine Core)/(Auto select Core) for external irq
- software irq send/broadcast any cores for IPI





linux/drivers/irqchip/irq-csky-mpintc.c:

IRQ handler flow:

```
static void csky_mpintc_handler(struct pt_regs *regs) {
    do {
        handle_domain_irq(root_domain,
                           readl_relaxed(reg_base + INTCL_RDYIR), regs);
    } while (readl_relaxed(reg_base + INTCL_HPPIR) & BIT(31));
}
```

```
static void csky_mpintc_eoi(struct irq_data *d) {
    writel_relaxed(d->hwirq, reg_base + INTCL_CACR);
}
```

---

IRQ enable/disable flow:

```
static void csky_mpintc_enable(struct irq_data *d) {
    writel_relaxed(d->hwirq, reg_base + INTCL_SENR);
}
```

```
static void csky_mpintc_disable(struct irq_data *d) {
    writel_relaxed(d->hwirq, reg_base + INTCL_CENR);
}
```

## IPI mechanism

```
static void csky_mpintc_send_ipi(const unsigned long *mask)
{
    /*
     * INTCL_SIGR[3:0] INTID
     * INTCL_SIGR[8:15] CPUMASK
     */
    writel_relaxed((*mask) << 8 | IPI_IRQ, reg_base + INTCL_SIGR);
}
```

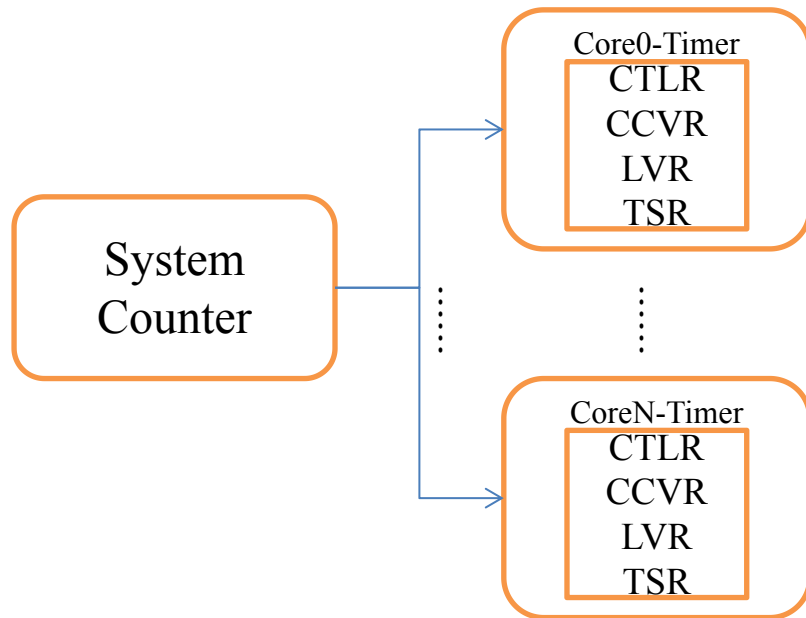
## IRQ TYPE Config

...

## IRQ Priority Config

...

- Private Interrupt Num: 16
- Count Control Register
- Count Current Value Register
- Load Value Register
- TS Clear Register



## linux/drivers/clocksource/timer-mp-csky.c

```
static irqreturn_t csky_timer_interrupt(int irq, void *dev) {
    struct timer_of *to = this_cpu_ptr(&csky_to);
    mtcrr(PTIM_TSR, 0);
    to->clkdev.event_handler(&to->clkdev);
    return IRQ_HANDLED;
}

static int csky_mptimer_set_next_event(unsigned long delta,
                                       struct clock_event_device *ce) {
    mtcrr(PTIM_LVR, delta);
    return 0;
}

static u64 clksrc_read(struct clocksource *c) {
    return (u64)mfcrr(PTIM_CCVR);
}
```



```
#define ATOMIC_OP(op, c_op)
static inline void atomic_##op(int i, atomic_t *v) {
    asm volatile (
        "1:      ldx.w          %0, (%2) \n"
        "        " #op          %0, %1  \n"
        "        stx.w          %0, (%2) \n"
        "        bez           %0, 1b   \n"
        : "=&r" (tmp)

```

- will add an entry into the local monitor, and the entry is composed of a address tag and a exclusive flag (initied with 1).
- Any stores (include other cores') will break the exclusive flag to 0 in the entry which could be indexed by the address tag.

- **Store Success:** When the entry's exclusive flag is 1, it will store `rz` to the `[rx + off]` address and the `rz` will be set to 1.
- **Store Failure:** When the entry's exclusive flag is 0, just `rz` will be set to 0.

© 2017 C-SKY MICROSYSTEMS CO., LTD All rights reserved.

```

/* Test-and-set spin-locking. */
arch_spin_lock(arch_spinlock_t *lock) {
"1:      ldex.w          %0, (%1) \n"
"      bnez             %0, 1b  \n"
"      movi             %0, 1  \n"
"      stex.w           %0, (%1) \n"
"      bez              %0, 1b  \n"
  smp_mb(); }

arch_spin_unlock(arch_spinlock_t *lock) {
  smp_mb();
"      movi             %0, 0  \n"
"      stw               %0, (%1) \n" }

```



/\* read lock/unlock/trylock \*/

```
arch_read_lock(arch_rwlock_t *lock) {
"1:      ldex.w      %0, (%1) \n"
"        blz        %0, 1b  \n"
"        addi       %0, 1  \n"
"        stex.w      %0, (%1) \n"
"        bez        %0, 1b  \n"
smp_mb(); }
```

```
arch_read_unlock(arch_rwlock_t *lock) {
smp_mb();
```

```
"1:      ldex.w      %0, (%1) \n"
"        subi       %0, 1  \n"
"        stex.w      %0, (%1) \n"
"        bez        %0, 1b  \n" }
```

/\* write lock/unlock/trylock \*/

```
arch_write_lock(arch_rwlock_t *lock) {
"1:      ldex.w      %0, (%1) \n"
"        bnez       %0, 1b  \n"
"        subi       %0, 1  \n"
"        stex.w      %0, (%1) \n"
"        bez        %0, 1b  \n"
smp_mb(); }
```

```
arch_write_unlock(arch_rwlock_t *lock) {
smp_mb();
```

```
"1:      ldex.w      %0, (%1) \n"
"        movi       %0, 0  \n"
"        stex.w      %0, (%1) \n"
"        bez        %0, 1b  \n" }
```

linux/arch/csky/include/asm/spinlock.h

```
arch_spin_lock(arch_spinlock_t *lock) {
    arch_spinlock_t lockval;
    u32 ticket_next = 1 << TICKET_NEXT; //16

    "1:          ldex.w          %0, (%2) \n"
    "          mov             %1, %0 \n"
    "          add              %0, %3 \n"
    "          stex.w           %0, (%2) \n"
    : "=&r" (tmp), "=&r" (lockval)
    : "r"(p), "r"(ticket_next)
    : "cc");

    while (lockval.tickets.next != lockval.tickets.owner)
        lockval.tickets.owner = READ_ONCE(lock->tickets.owner);
        smp_mb();
    }

    arch_spin_unlock(arch_spinlock_t *lock) {
        smp_mb();
        lock->tickets.owner++;
    }
}
```





Documentation/memory-barriers.txt:

For example, consider the following sequence of events:

CPU 1	CPU 2
=====	=====
{ A == 1; B == 2 }	
A = 3;	x = B;
B = 4;	y = A;

The set of accesses as seen by the memory system in the middle can be arranged in 24 different combinations:

STORE A=3,	STORE B=4,	y=LOAD A->3,	x=LOAD B->4
STORE A=3,	STORE B=4,	x=LOAD B->4,	y=LOAD A->3
STORE A=3,	y=LOAD A->3,	STORE B=4,	x=LOAD B->4
STORE A=3,	y=LOAD A->3,	x=LOAD B->2,	STORE B=4
STORE A=3,	x=LOAD B->2,	STORE B=4,	y=LOAD A->3
STORE A=3,	x=LOAD B->2,	y=LOAD A->3,	STORE B=4
STORE B=4,	STORE A=3,	y=LOAD A->3,	x=LOAD B->4
STORE B=4,	...		
...			

sync: completion barrier  
sync. s: completion barrier and shareable to other cores  
sync. i: completion barrier with flush cpu pipeline  
sync. is: completion barrier with flush cpu pipeline and shareable to other cores

bar. brwarw: ordering barrier for all load/store instructions before it  
bar. brwarws: ordering barrier for all load/store instructions before it and shareable to other cores  
bar. brar: ordering barrier for all load instructions before it  
bar. brars: ordering barrier for all load instructions before it and shareable to other cores  
bar. bwaw: ordering barrier for all store instructions before it  
bar. bwaws: ordering barrier for all store instructions before it and shareable to other cores

arch/csky/mm/tlb.c:

- `tlbi.va(s) rx` Delete va-tlb-entry with care ASID
- `tlbi.vaa(s) rx` Delete va-tlb-entry without care ASID
- `tlbi.asid(s) rx` Delete all the ASID tlb-entries
- `tlbi.all(s)` Delete all the ASID tlb-entries

arch/csky/mm/cachev2.c: (shareable in tlb-entry)

- `icache.iva rx` invalid va in icache
- `dcache.iva rx` invalid va in dcache
- `dcache.cval1 rx` flush va in dcache
- `dcache.cva rx` flush va in dcache & l2cache
- `dcache.civa rx` flush and invalid va in dcache & l2cache

Guide: <https://c-sky.github.io/>

```
$ xz -d vmlinux.xz
$ mkdir toolchain
$ cd toolchain
$ tar -Jxf ../qemu_csky_ck860_4.18_glibc_defconfig.tar.xz
$ cd ..
$ toolchain/csky-qemu/bin/qemu-system-csky2 -kernel vmlinux -dtb qemu.dtb
-nographic -M mp860 -smp 4
```

# Thank You

# Any Question is welcome!

[www.c-sky.com](http://www.c-sky.com)

