

Toward General Deadlock of Locks Prediction in Lockdep

Du, Yuyang

Outline

- Introduce deadlock
- Demystify lockdep
- Read-write locks
- General deadlock of locks prediction algorithm

Deadlock



(图片来自网络)

Every car is alive but cannot
make forward progress forever!

This happens on multi-threaded or distributed-system programs too.

Why deadlock happens?



(图片来自网络)

The story of three monks no water.

A deadlock is caused by circular waiting relationship among the actors.

Deadlocks are bad


- If happened, it is devastating
- Concrete vs. potential deadlocks
 - E.g., timing conditions are met
- If not so far, it will sometime according to Murphy's Law
 - "Anything that can go wrong will go wrong"
- Deadlock is notoriously hard to debug

What can be done about deadlocks?

- Detection
 - After concrete deadlocks happened
- Prevention
 - Before concrete deadlocks happen
- Prediction
 - Find potential deadlocks
 - A mechanism to characterize the locking behavior of workloads
 - Then based on the behavior report potential deadlocks

Deadlock of locks

Example:



What a long circle! So complicated!

Waitings:

Thread1 Lock B -> Thread2 Unlock B
Thread2 Unlock B -> Thread2 Unlock A
Thread2 Unlock A -> Thread2 Lock A
Thread2 Lock A -> Thread1 Unlock A
Thread1 Unlock A -> Thread1 Unlock B
Thread1 Unlock B -> Thread1 Lock B

What locking behavior to characterize

- When deadlock happens, threads are attempting to acquire locks while holding another

Thread

Lock A

Lock B

Lock C

...

How do we describe such lock holding-and-attempting's?

Lock dependency – locking order

- Whenever a thread attempts to acquire lock B while hold lock A

Thread

Lock A

Lock B

Unlock B

Unlock A

- There exists the locking order

A -> B

Use lock dependency to find deadlocks

The previous example:

Thread1

Thread2

Lock A

Lock B

Lock B

Lock A



characterize

A -> B



characterize

B -> A



Potential deadlock!

Lock dependency vs. waiting relationship

- Whenever a thread attempts to acquire lock B while holding lock A

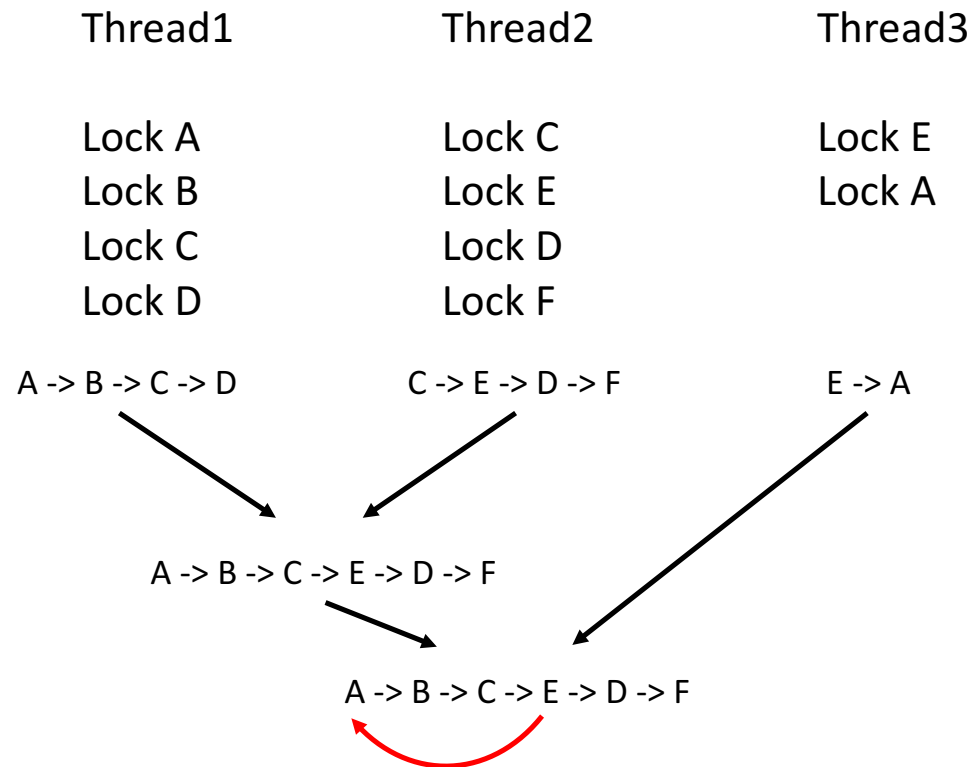
Thread1	Threadx
Lock A	
Lock B	Lock A
Unlock B	
Unlock A	

- It is very probable there is a Threadx having Lock A, then:

Threadx Lock A -> Thread1 Unlock A	}	Threadx Lock A -> Thread1 Lock B	}	A -> B
Thread1 Unlock A -> Thread1 Unlock B				
Thread1 Unlock B -> Thread1 Lock B				

Lock dependency in general

- Whenever a thread attempts to acquire a lock while hold some locks



That is it!

- You have known almost all about what lockdep is doing
- Except
 - What if interrupts come into play?
 - Safe and unsafe lock.
 - Why can it predict?
 - Lock classification
 - Lock waiting pattern
 - Why is there false positive?
 - Incomplete relative timings
 - Exaggerated classification
 - How can it improve performance?
 - Lock chain caching, etc.

Except: when read-write locks come into play

- Read locks (readers) and write locks (writers)
- Exclusive lock vs. read-write lock
 - Write lock == exclusive lock
 - Read locks (only) can be concurrent
- Read-write locks are the general-form locks

But the previous deadlock detection scheme only works for exclusive locks. Why?

Read-write lock design elements

- Can readers be recursive?

Thread

RLock A

RLock A

Unlock A

Unlock A

- Lock grant order when readers and writers contend

rwsem

- No recursive
- Grant order: prioritize writers

Notation:

Rn: Read lock n

Wn: Write lock n

RRn: Recursive-read lock n

Xn: Read, recursive-read or
write lock n



rwlock

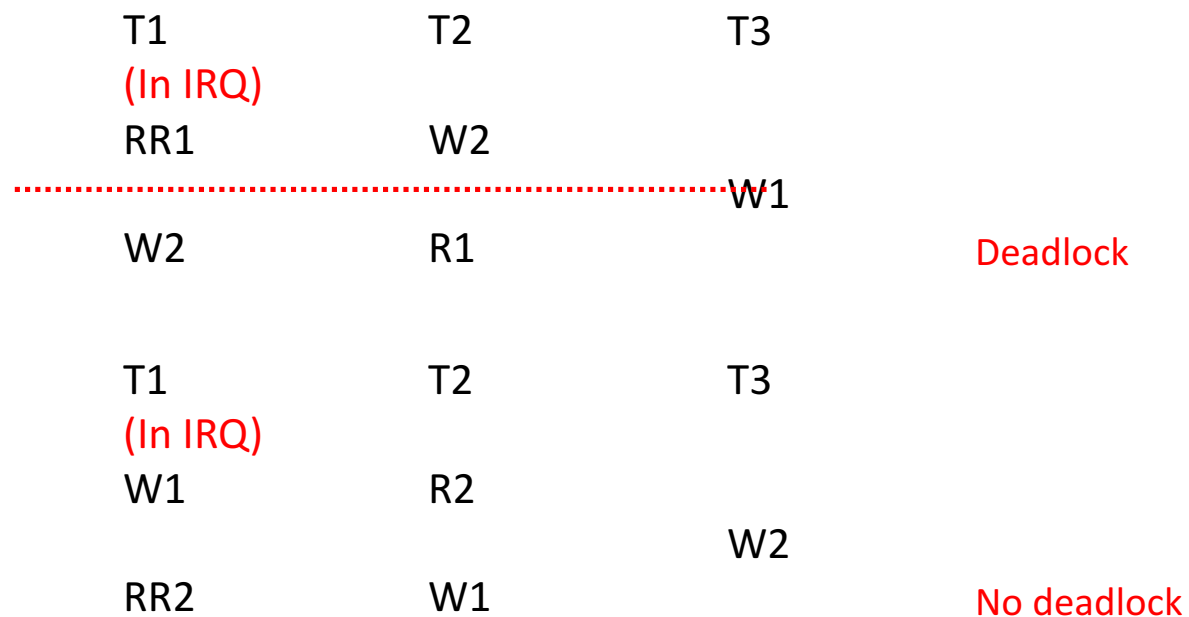
- Recursive
- Grant order: prioritize readers

T1	T2	
RR1	W2	
W2	RR1	No deadlock

T1	T2	T3	
RR1	W2		
W2	RR1	W1	
			No deadlock

qrwlock – an implementation of rwlock

- Recursive
- Grant order: readers or writers priority depends on contexts



The way to general deadlock prediction algorithm

- 7 lemmas lead to (and proves) the solution
- Using abstract cases to show all possibilities
- Step 1: propose the Simple Algorithm
- Step 2: adjust the Simple Algorithm to the Final Algorithm

Lemma1: what makes a deadlock prediction?

- Circular waiting relationship is **necessary** for deadlocks
- But they are **not sufficient** with (recursive) read locks

Lemma #1

A deadlock can be and can only be detected at the earliest time at the final arc that completes a waiting circle. In other words, at the final arc, the problem is try to find out whether this circle to come is a deadlock or not.

Lemma2: number of threads in deadlock

- Given a deadlock, assume there are n threads, denoted as:

$$T^1, T^2, \dots, T^n$$

- Depending on n , there are two cases:
 - $n = 1$, referred to as recursion deadlock scenario, skipped
 - $n \geq 2$, grouped as (T^1, \dots, T^{n-1}) and T^n . By consolidating the former, we get a big imagined $T1$ and $T2$ (with task numbers adjusted).

Lemma #2

Two threads ($T1$ and $T2$) can virtually represent any situation with any number of tasks to check for deadlock.

Locking behavior representation

- Two-thread model (version 1)
 - T1 - the entire previous locking behavior depicted by the lock dependency graph
 - T2 - the current new locking behavior, the X1 -> X2 dependency
- Caveats

Case #1.1

T1	T2
W1	
RR2	W3
W3	W1

Deadlock

Case #1.2

T1	T2
W1	
RR2	
	W3
RR2	W1
W3	

No deadlock

Case #1.3

T1a	T1b	T2
W1	RR2	W3
RR2	W3	W1


No deadlock

Lemma3: numbers of locks do not matter!

- With the two-thread model, the number of threads does not matter
- Further, the number of locks does not matter either


Lemma #3

Any deadlock scenario can be converted to an ABBA deadlock.



General problem to solve

- Given the two-thread model: T1 and T2
- When T2 has a BA to come, is there an AB in T1?



Lock type exclusiveness table

- Lock type
 - Read, recursive-read, or write
- Is A exclusive to B?

T1

T2

or

T1

T2

A

B

B

A

A vs. B	Recursive-read lock	Read lock	Write lock
Recursive-read lock	No	Yes	Yes
Read lock	No	Yes	Yes
Write lock	Yes	Yes	Yes

Simple Algorithm

- Given T2's dependency, divide and conquer the cases respectively
- There are 9 cases
 - Read lock and read lock
 - Read lock and recursive-read lock
 - Read lock and write lock
 - Write lock and read lock
 - Write lock and recursive-read lock
 - Write lock and write lock
 - Recursive-read lock and read lock
 - Recursive-read lock and recursive-read lock
 - Recursive-read lock and write lock

Simple Algorithm case

- When the final dependency is ended with read lock and read lock

Case #2.1

T1	T2
X1	R2
W2	R1

Deadlock

Case #2.2

T1	T2
X1	R2
R2	R1

Deadlock

Case #2.3

T1	T2
X1	R2
RR2	R1

No deadlock

Simple Algorithm cont.


- Given T1 and T2, and an ABBA locks

T1	T2
X1.A	X2.A
X2.B	X1.B

- Simple algorithm
 - (1) If X1.A vs. X1.B are exclusive and X2.A vs. X2.B are exclusive then it is deadlock
 - (2) Otherwise no deadlock

Lock type promotion

- In locking Lock type in dependency is promoted toward more exclusiveness



Final Algorithm (1/6)

- A case fails the Simple Algorithm

Case #13

T1	T2
X1	
RR2	RR2
X3	X1

No deadlock

Case #12

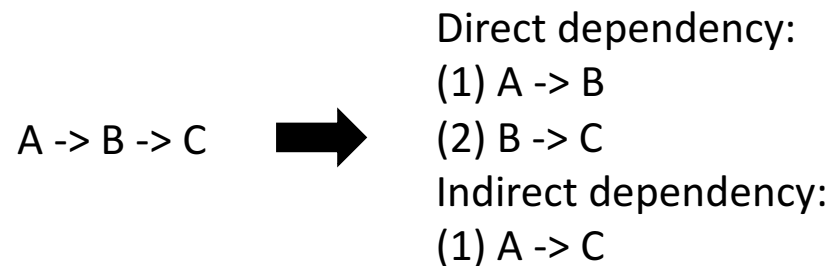
T1	T2
X1	X3
RR2	RR2
X3	X1

Deadlock

Indirect dependency: X3 -> X1

Final Algorithm (2/6)

- Direct dependency vs. indirect dependency



Lemma #4

The direct dependency (e.g., $RR2 \rightarrow X1$ in Case #12) that serves in the path from $X3 \rightarrow X1$ is **critical**.

Final Algorithm (3/6)

- Case #12 vs. Case #13

Lemma #5

Missed in Case #13, the game changer to Case #12 is that it has X3 in T2 whereas Case #13 does not.

- From the Simple Algorithm to the Final Algorithm
 - (a) Continue searching the graph to find a new circle.
 - (b) In the new circle, do the Simple Algorithm, if a deadlock is found, done
 - (c) In the new circle, if previous locks in T2's stack (or chain) are in it and then check whether the circle is indeed a deadlock. This is done by checking each newly introduced indirect dependency (such as X3 - > X1 in Case #12) according to our Simple Algorithm as before.
 - (d) If a deadlock is found then the algorithm is done, otherwise go to (a) unless the entire graph is traversed

Final Algorithm (4/6)

- The failed case can be solved now

Case #12

T1	T2
X1	X3
RR2	RR2
X3	X1

T1
Previous dependency graph



T2
New dependency 2 -> 1

First: Do we have 1 -> 2 ? Yes
Second: do we have 1 -> 3? Yes.

Final Algorithm (5/6)

- Why does the Final Algorithm work?

Lemma #6

Lemma #5 nicely raises a question whether a previous lock involved indirect dependency in T2 is **necessary**. The answer is yes, otherwise our **Simple Algorithm** has already solved the problem.

- Modified two-thread model (version 2)
 - T1 - the entire previous locking behavior depicted by the lock dependency graph.
 - T2 - the current task's held lock stack worth of direct and indirect dependencies.

Final algorithm (6/6)

- Why does the final algorithm work?

Lemma #7

It is also natural to ask whether indirect dependencies with a starting lock in T2 only is **sufficient**: what if the indirect dependency has a starting lock from T1. The answer is yes too.

- More explanation
 - Because Lemma #2 and Lemma #3 say that any deadlock is an ABBA so that T1 can only contribute an AB and T2 must have a BA, and the final direct arc or dependency is the BA or the ending part of the BA.

Questions?

*Thanks,
Yuyang Du*

Backup

- Backup